

# *DS-GA 3001.009: Responsible Data Science*

## Data Profiling continued

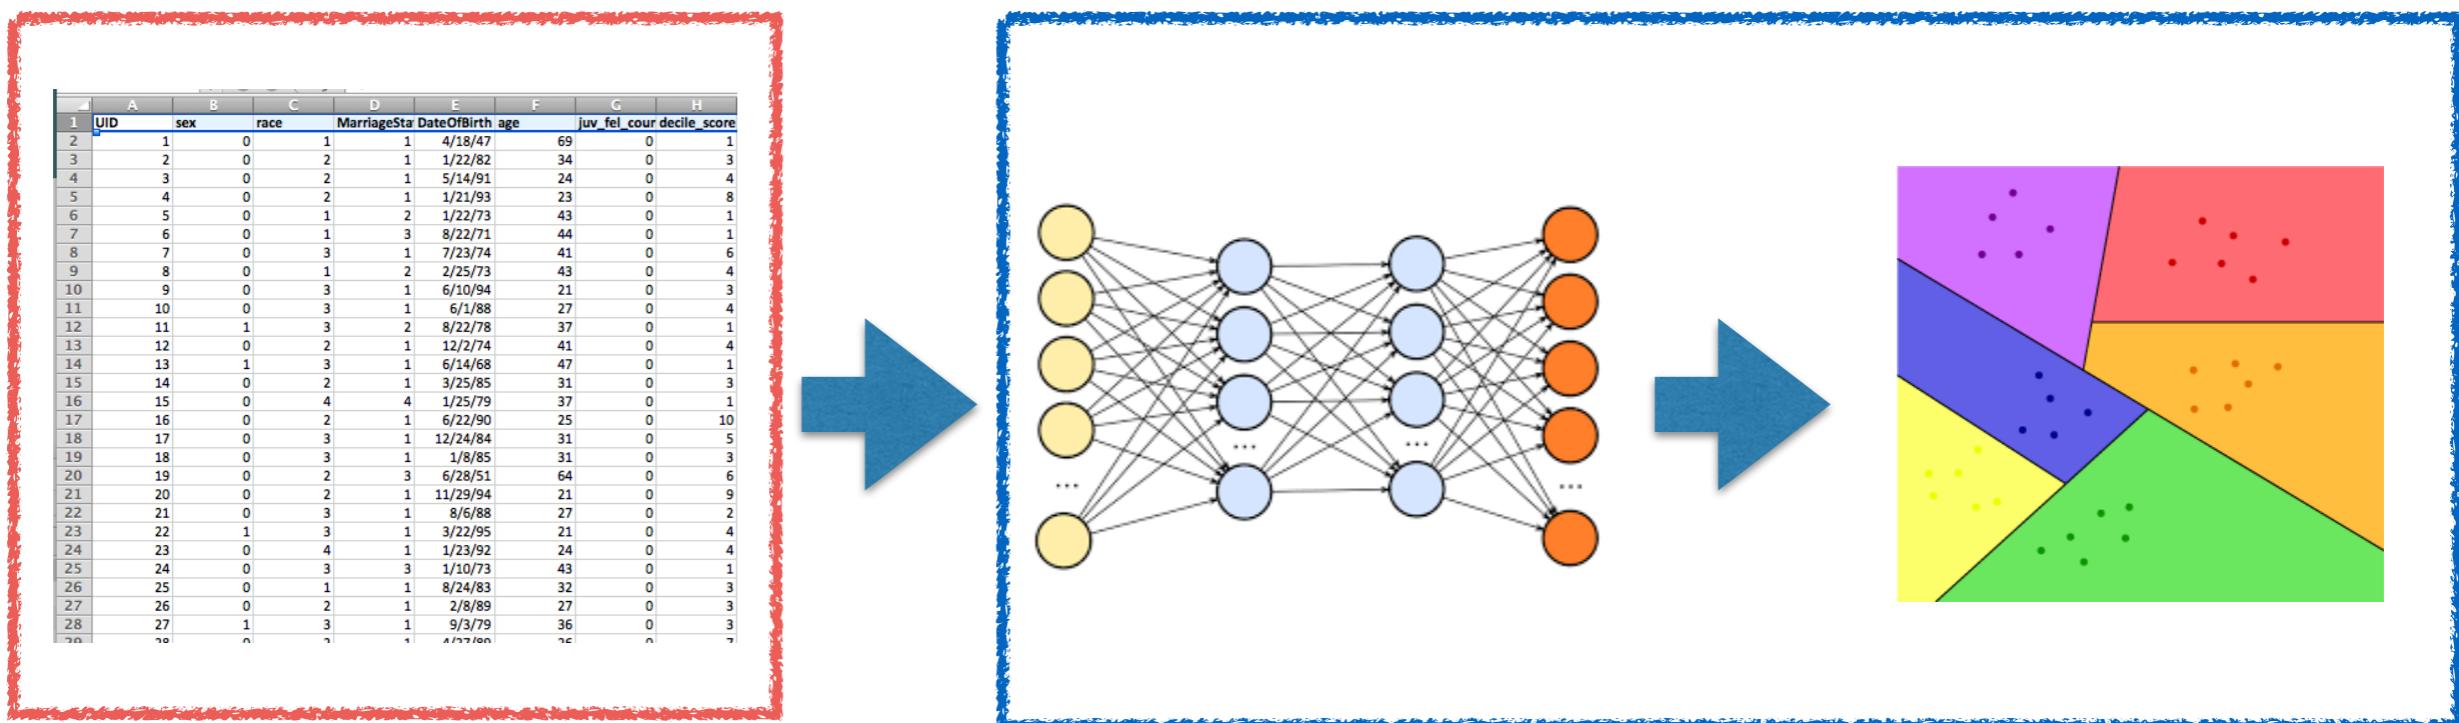
Prof. Julia Stoyanovich  
Center for Data Science  
Computer Science and Engineering at Tandon

@stoyanoj

<http://stoyanovich.org/>  
<https://dataresponsibly.github.io/>

# Responsible data science

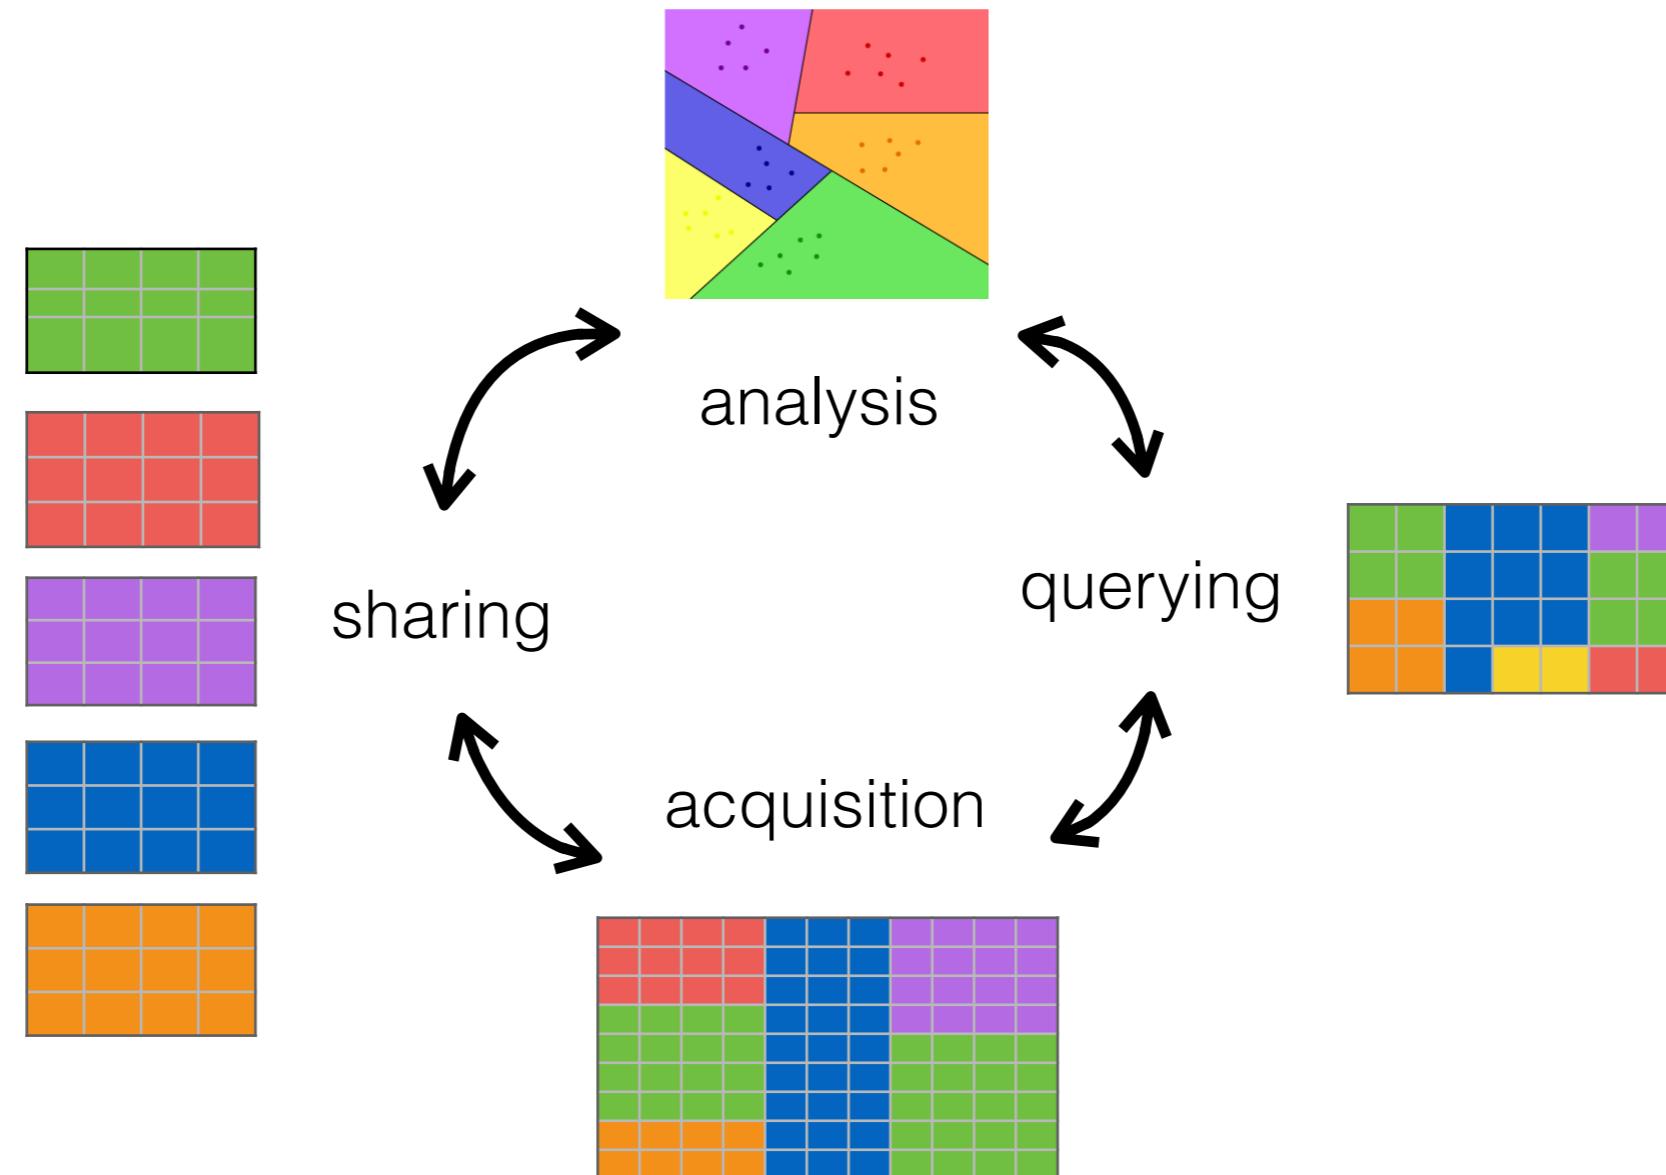
- Be **transparent** and **accountable**
- Achieve **equitable** resource distribution
- Be cognizant of the **rights** and **preferences** of individuals



done?

but where does the data come from?

# A holistic view of the lifecycle



# Understand your data!



“Given the heterogeneity of the flood of data, it is **not enough merely to record it and throw it into a repository**. Consider, for example, data from a range of scientific experiments. If we just have a bunch of data sets in a repository, it is **unlikely anyone will ever be able to find, let alone reuse**, any of this data. With adequate **metadata**, there is some hope, but even so, challenges will remain due to differences in experimental details and in data record structure.”

<https://cra.org/ccc/wp-content/uploads/sites/2/2015/05/bigdatawhitepaper.pdf>

# Understand your data!

## 2.2 Big data



In the analog age, most of the data that were used for social research was created for the purpose of doing research. In the digital age, however, a huge amount of

**data is being created by companies and governments for purposes other than research,**

such as providing services, generating profit, and administering laws. Creative people, however, have realized that you can **repurpose** this corporate and government data for research.

<https://www.bitbybitbook.com/en/1st-ed/observing-behavior/data/>

# Understand your data!

## 2.2 Big data



... from the perspective of researchers, big data sources are “found,” they don’t just fall from the sky. Instead, data sources that are “found” by researchers are **designed by someone for some purpose**. Because “found” data are designed by someone, I always recommend that you **try to understand as much as possible about the people and processes that created your data**.

<https://www.bitbybitbook.com/en/1st-ed/observing-behavior/data/>

# Understand your data!

Need **metadata** to:

- enable data **re-use** (have to be able to find it!)
- determine **fitness for use** of a dataset in a task
- help establish **trust** in the data analysis process and its outcomes

Data is considered to be of high quality if it's "**fit for intended uses** in operations, decision making and planning"

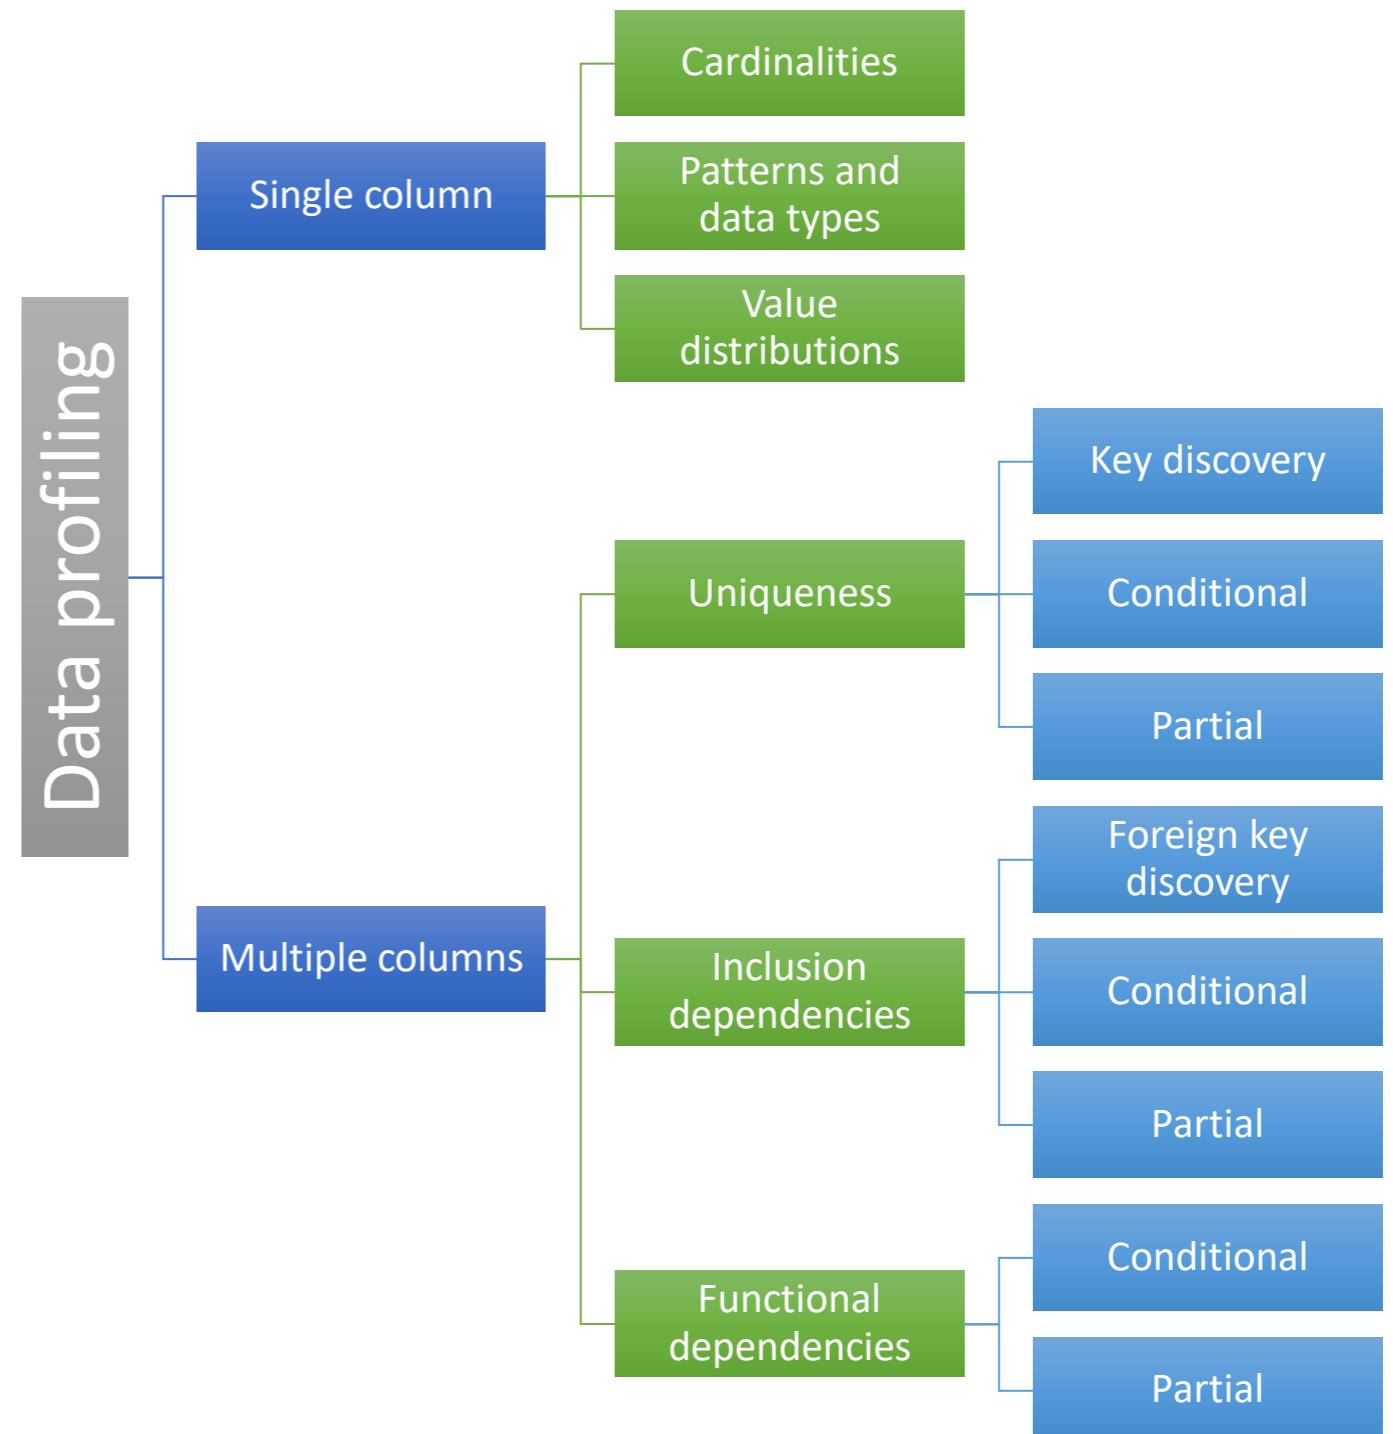
[Thomas C. Redman, "Data Driven: Profiting from Your Most Important Business Asset." 2013]

# A classification of data profiling tasks

[Abedjan, Golab, Naumann; *SIGMOD 2017*]

	A	B	C	D	E	F	G	H
1	UID	sex	race	MarriageSta	DateOfBirth	age	juv_fel_cour	decile_score
2	1	0	1	1	4/18/47	69	0	1
3	2	0	2	1	1/22/82	34	0	3
4	3	0	2	1	5/14/91	24	0	4
5	4	0	2	1	1/21/93	23	0	8
6	5	0	1	2	1/22/73	43	0	1
7	6	0	1	3	8/22/71	44	0	1
8	7	0	3	1	7/23/74	41	0	6
9	8	0	1	2	2/25/73	43	0	4
10	9	0	3	1	6/10/94	21	0	3
11	10	0	3	1	6/1/88	27	0	4
12	11	1	3	2	8/22/78	37	0	1
13	12	0	2	1	12/2/74	41	0	4
14	13	1	3	1	6/14/68	47	0	1
15	14	0	2	1	3/25/85	31	0	3
16	15	0	4	4	1/25/79	37	0	1
17	16	0	2	1	6/22/90	25	0	10
18	17	0	3	1	12/24/84	31	0	5
19	18	0	3	1	1/8/85	31	0	3
20	19	0	2	3	6/28/51	64	0	6
21	20	0	2	1	11/29/94	21	0	9
22	21	0	3	1	8/6/88	27	0	2
23	22	1	3	1	3/22/95	21	0	4
24	23	0	4	1	1/23/92	24	0	4
25	24	0	3	3	1/10/73	43	0	1
26	25	0	1	1	8/24/83	32	0	3
27	26	0	2	1	2/8/89	27	0	3
28	27	1	3	1	9/3/79	36	0	3
29	28	0	3	1	4/17/00	26	0	7

relational data (here: just one table)



# An alternative classification

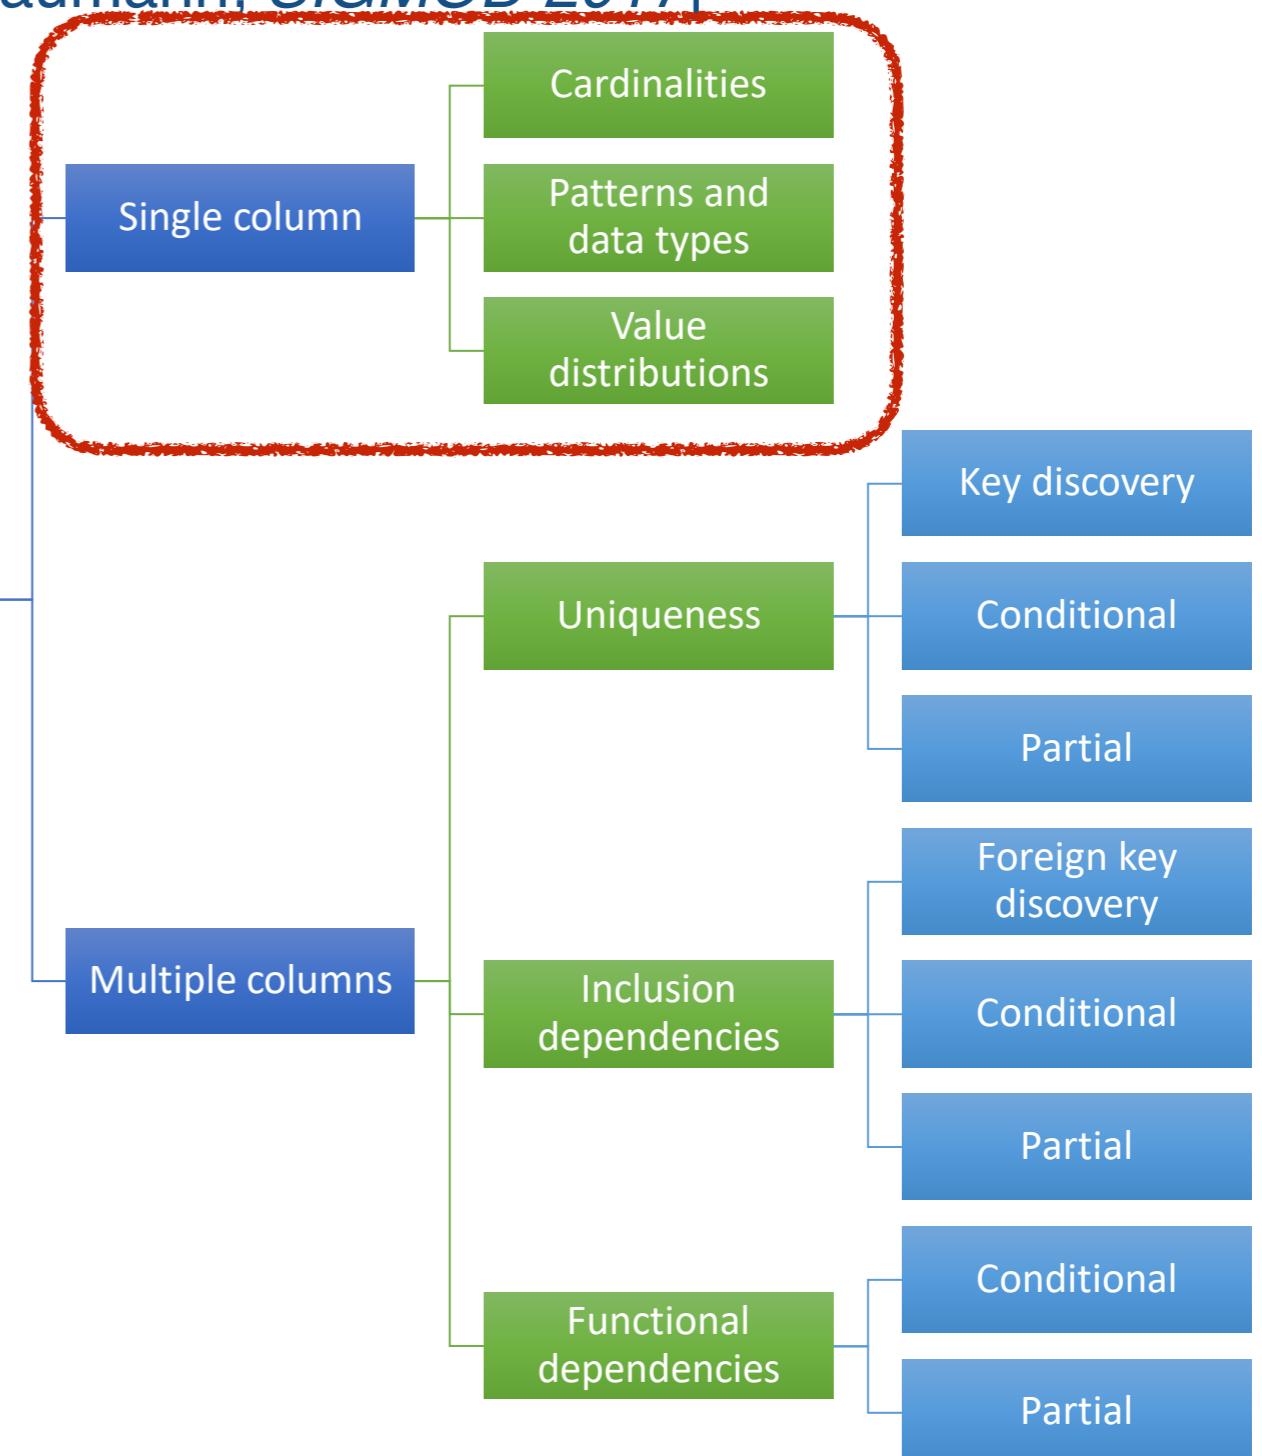
- To help understand the **statistics**, we look at value ranges, data types, value distributions per column or across columns, etc
- To help understand the **structure** - the (business) rules that generated the data - we look at unique columns / column combinations, dependencies between columns, etc - **reverse-engineer the relational schema** of the data we have
- We need both statistics and structure, they are mutually-reinforcing, and help us understand the **semantics** of the data - it's meaning

# A classification of data profiling tasks

[Abedjan, Golab, Naumann; *SIGMOD 2017*]

A	B	C	D	E	F	G	H	
1	UID	sex	race	MarriageSta	DateOfBirth	age	juv_fel_cour	decile_score
2	1	0	1	1	4/18/47	69	0	1
3	2	0	2	1	1/22/82	34	0	3
4	3	0	2	1	5/14/91	24	0	4
5	4	0	2	1	1/21/93	23	0	8
6	5	0	1	2	1/22/73	43	0	1
7	6	0	1	3	8/22/71	44	0	1
8	7	0	3	1	7/23/74	41	0	6
9	8	0	1	2	2/25/73	43	0	4
10	9	0	3	1	6/10/94	21	0	3
11	10	0	3	1	6/1/88	27	0	4
12	11	1	3	2	8/22/78	37	0	1
13	12	0	2	1	12/2/74	41	0	4
14	13	1	3	1	6/14/68	47	0	1
15	14	0	2	1	3/25/85	31	0	3
16	15	0	4	4	1/25/79	37	0	1
17	16	0	2	1	6/22/90	25	0	10
18	17	0	3	1	12/24/84	31	0	5
19	18	0	3	1	1/8/85	31	0	3
20	19	0	2	3	6/28/51	64	0	6
21	20	0	2	1	11/29/94	21	0	9
22	21	0	3	1	8/6/88	27	0	2
23	22	1	3	1	3/22/95	21	0	4
24	23	0	4	1	1/23/92	24	0	4
25	24	0	3	3	1/10/73	43	0	1
26	25	0	1	1	8/24/83	32	0	3
27	26	0	2	1	2/8/89	27	0	3
28	27	1	3	1	9/3/79	36	0	3
29	28	0	3	1	4/17/00	26	0	7

Data profiling



relational data (here: just one table)

# Single column: cardinalities, data types

[Abedjan, Golab, Naumann; *SIGMOD 2017*]

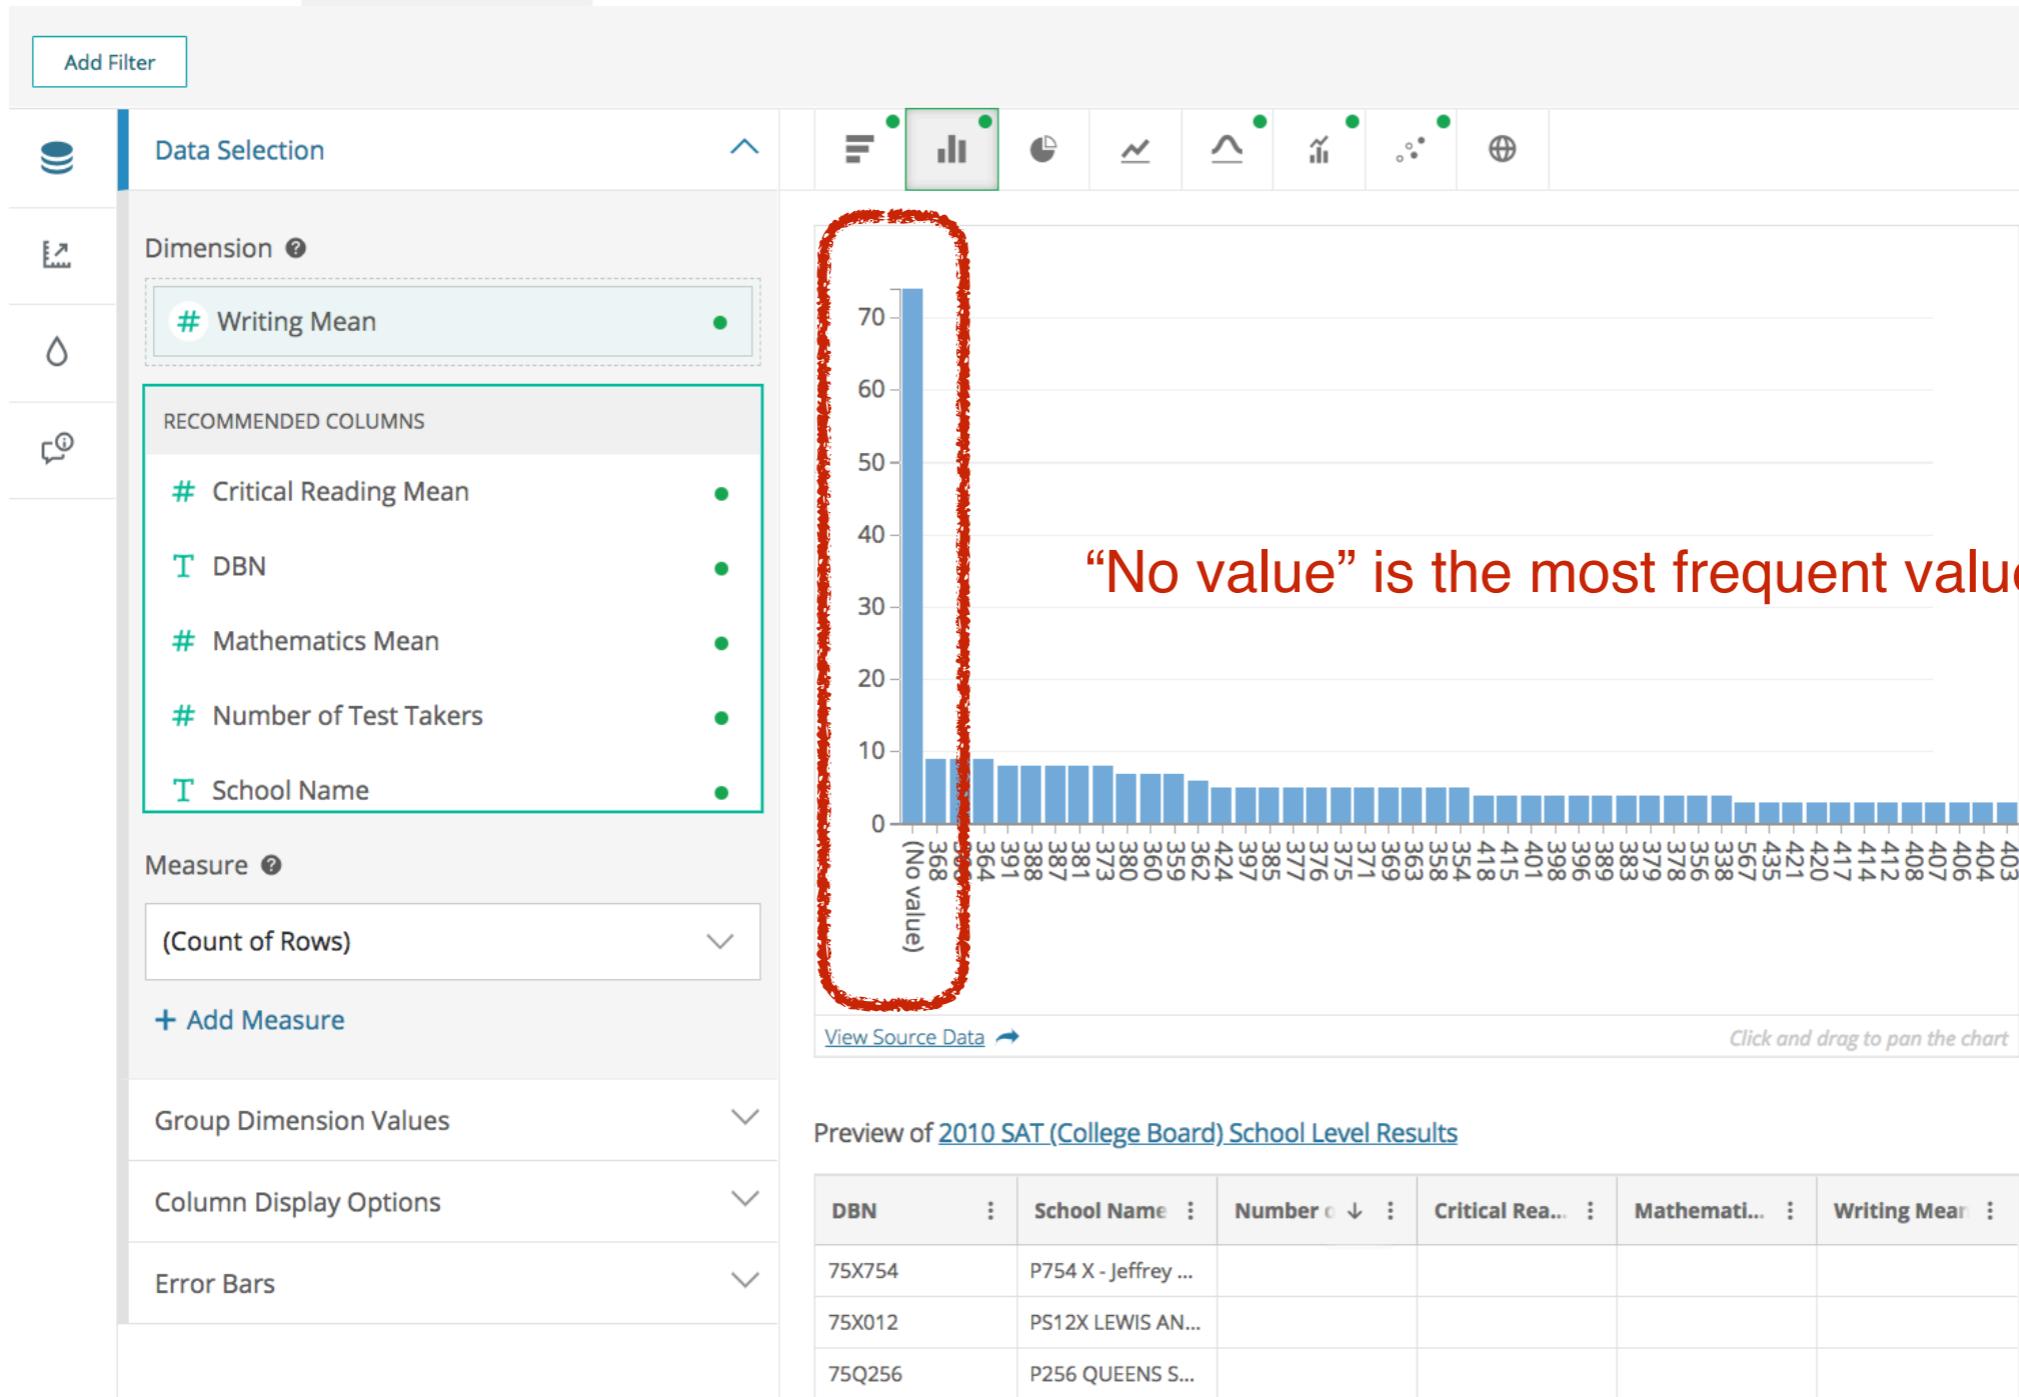
- cardinality of relation **R** - number of rows
- domain cardinality of a column **R.a** - number of **distinct** values
- attribute value **length**: min, max, average, median
- **basic data type**: string, numeric, date, time, ....
- number of percentage of **null** values of a given attribute
- regular expressions
- semantic domain: SSN, phone number
- ....

# 2010 SAT (College Board) School Level

## Results

Education

NYC OpenData



# 50 shades of null

- **Unknown** - some value definitely belongs here, but I don't know what it is (e.g., unknown birthdate)
- **Inapplicable** - no value makes sense here (e.g., if marital status = single then spouse name should not have a value)
- **Unintentionally omitted** - values is left unspecified unintentionally, by mistake
- **Optional** - a value may legitimately be left unspecified (e.g., middle name)
- **Intentionally withheld** (e.g., an unlisted phone number)
- .....

(this selection is mine, see reference below for a slightly different list)

<https://www.vertabelo.com/blog/technical-articles/50-shades-of-null-or-how-a-billion-dollar-mistake-has-been-stalking-a-whole-industry-for-decades>

# 50 shades of null... and it gets worse!

- **Hidden missing values** -
  - 99999 for zip code, Alabama for state
  - need data cleaning....
- lots of houses in Philadelphia, PA were built in 1934 (or 1936?) - not really!

how do we detect hidden missing values?

# Single column: cardinalities, data types

[Abedjan, Golab, Naumann; *SIGMOD 2017*]

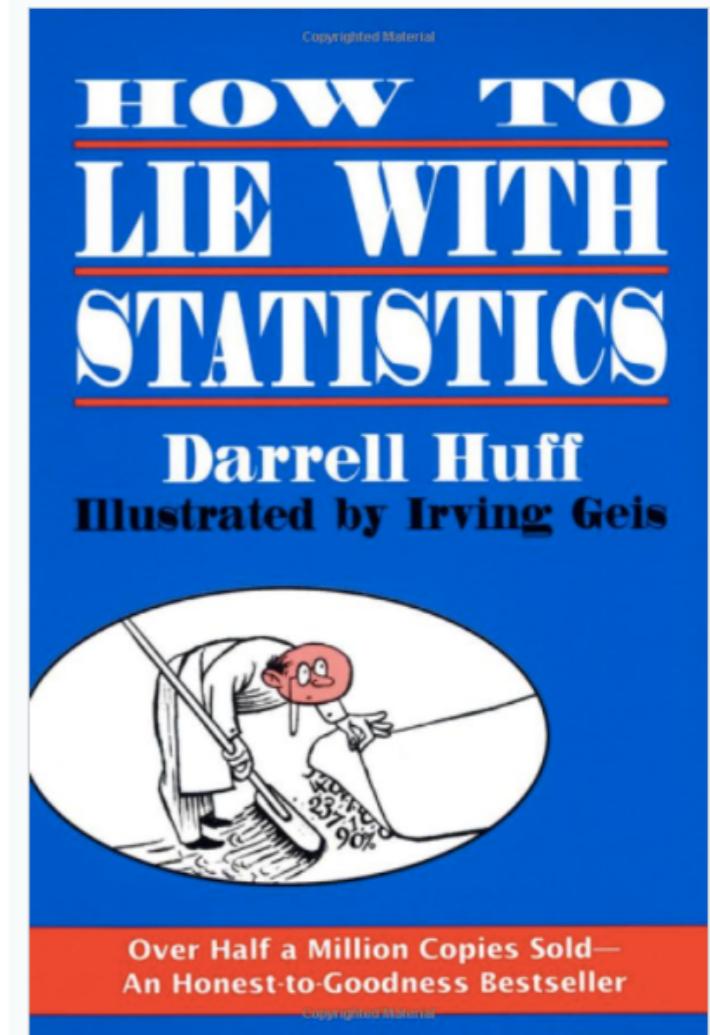
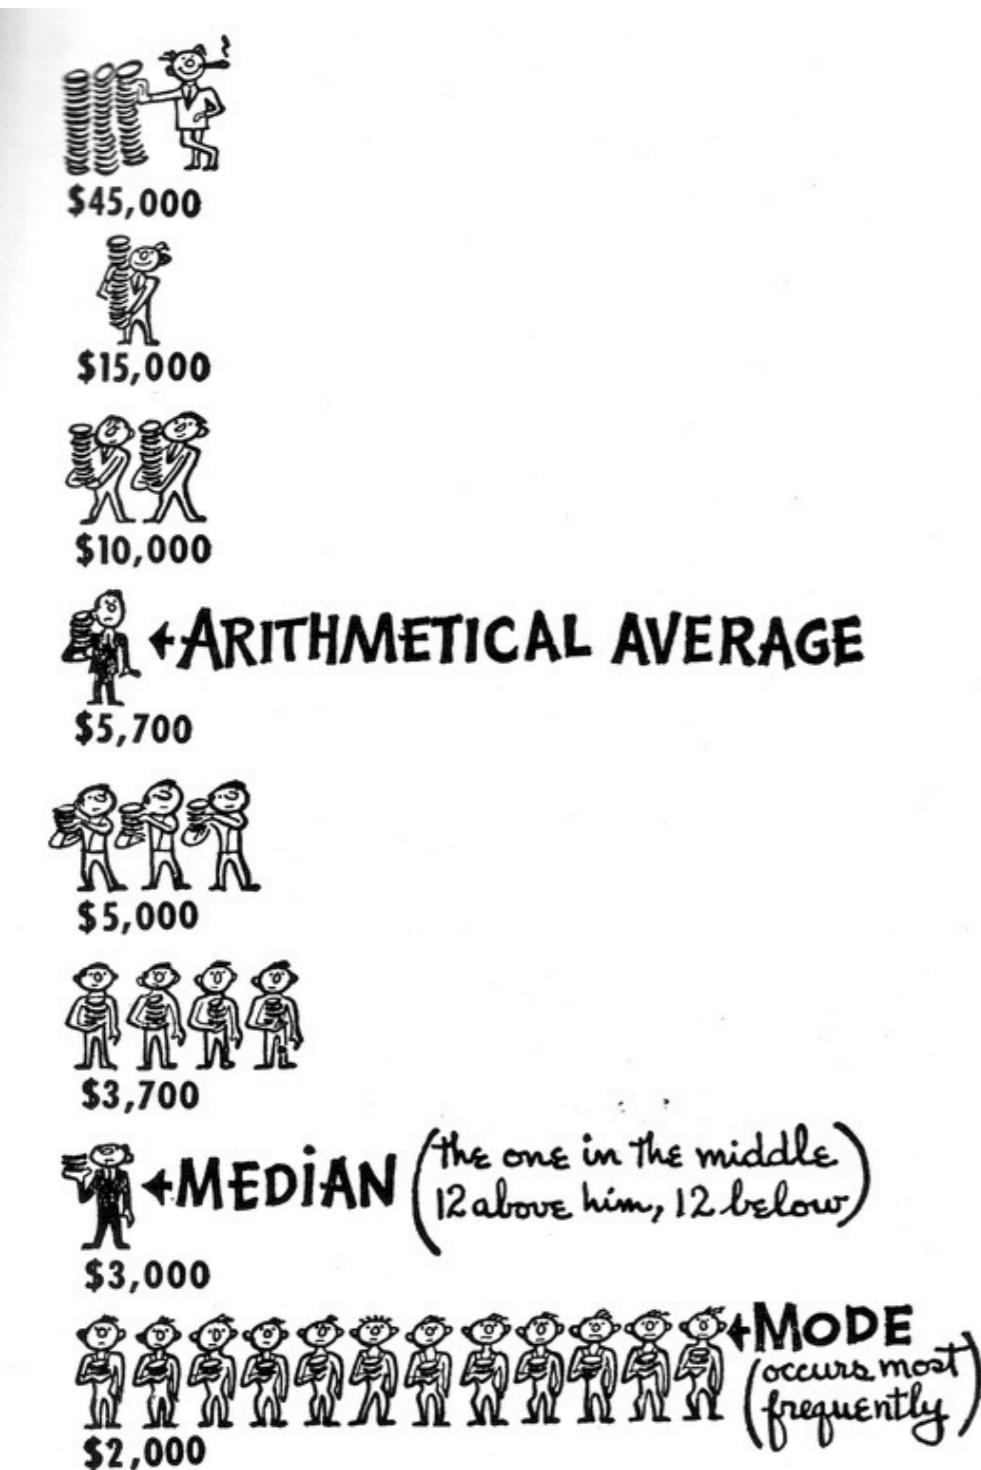
- cardinality of relation **R** - number of rows
- domain cardinality of a column **R.a** - number of **distinct** values
- attribute value **length**: min, max, average, median
- **basic data type**: string, numeric, date, time, ....
- number of percentage of **null** values of a given attribute
- **regular expressions**
- semantic domain: SSN, phone number
- ....

# Single column: basic stats, distributions

[Abedjan, Golab, Naumann; *SIGMOD 2017*]

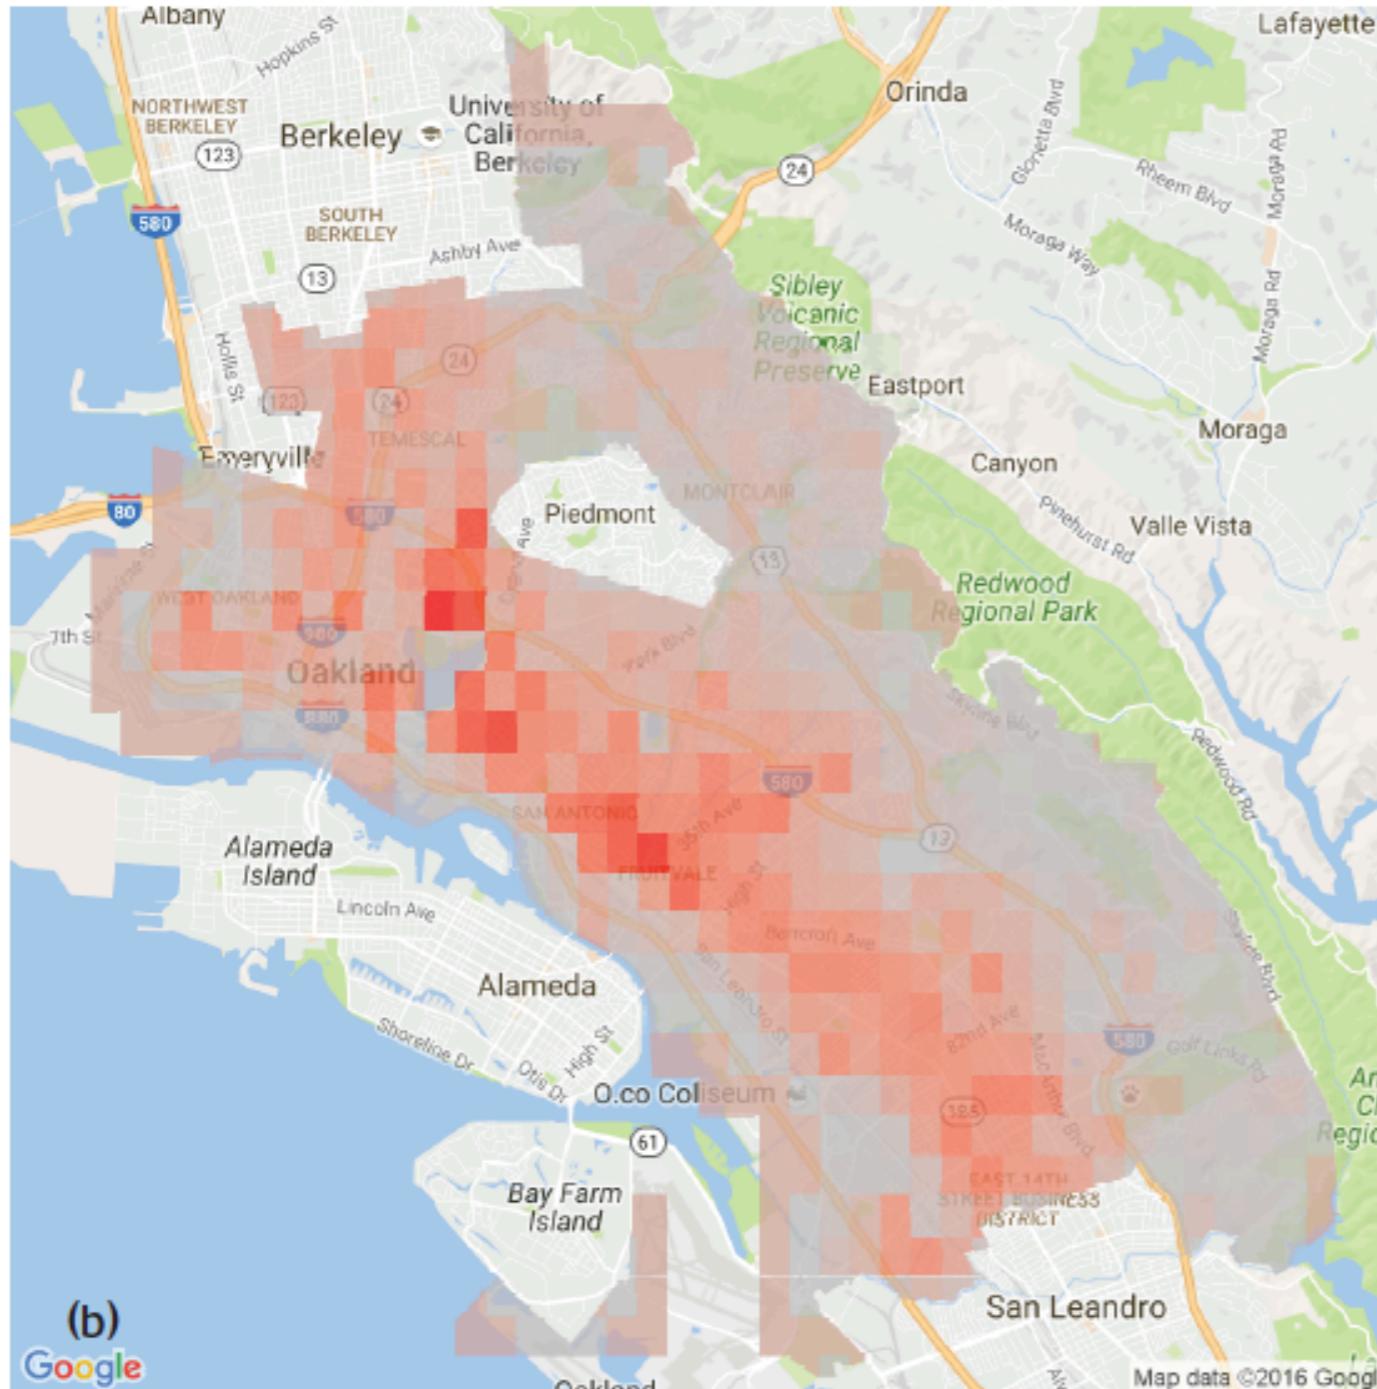
- min, max, **average**, median value of **R.a**
- **histogram**
  - equi-width - (approximately) the same number of distinct values in each bucket (e.g., age broken down into 5-year windows)
  - equi-depth (approximately) the same number of tuples in each bucket
  - biased histograms use different granularities for different parts of the value range to provide better accuracy
- quartiles - three points that divide the numeric values into four equal groups - a kind of an equi-depth histogram
- **first digit** - distribution of first digit in numeric values, to check Benford law
- ...

# The well-chosen average

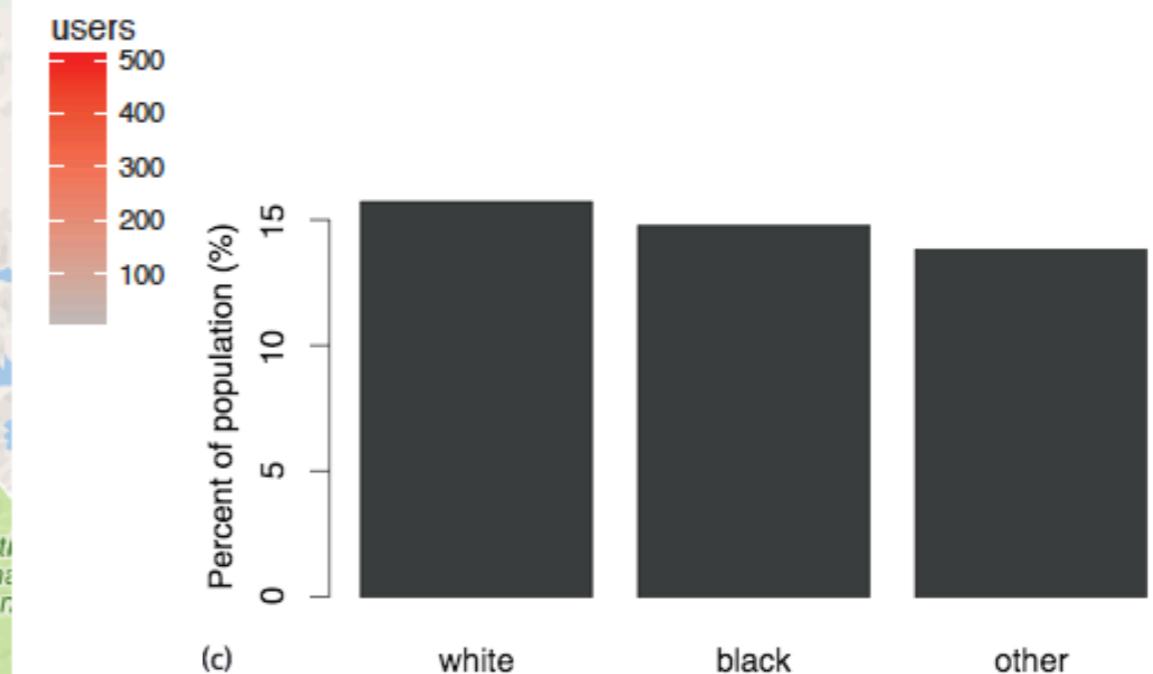


# Is my data biased? (histograms + geo)

[Lum, Isaac; *Significance*, 2016]



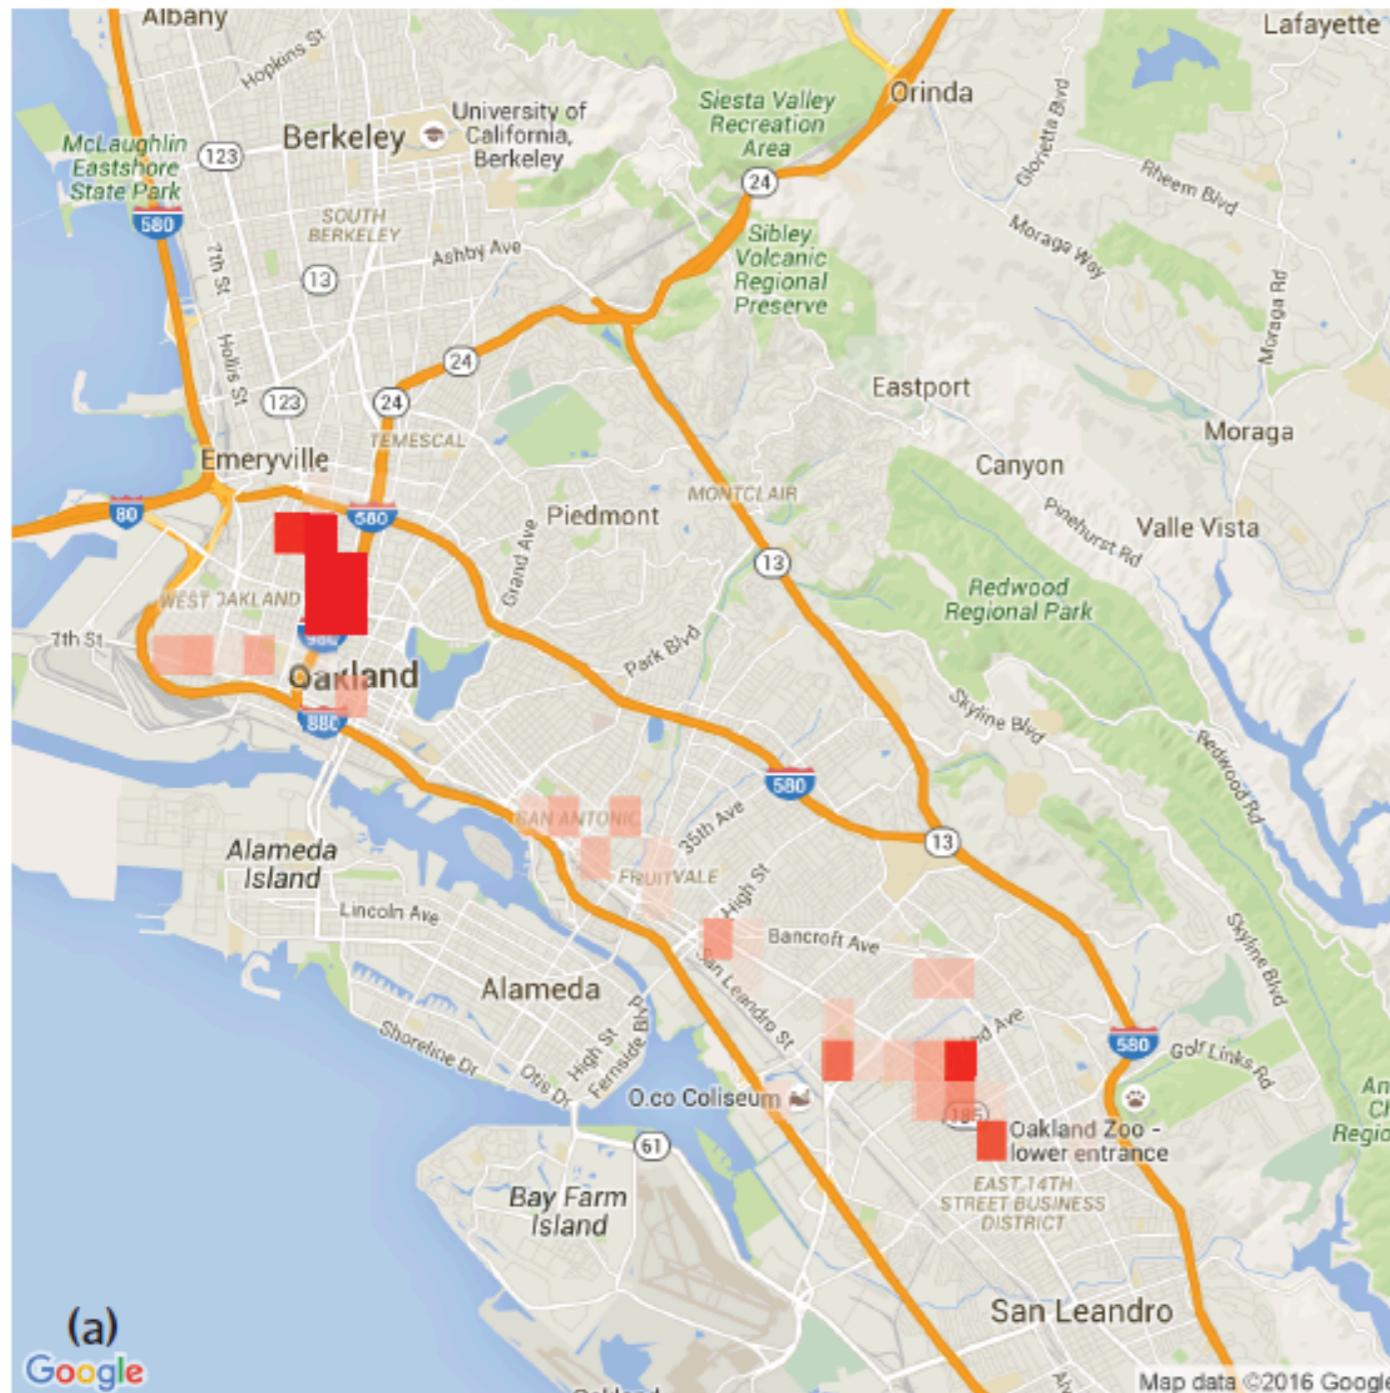
Estimated number of drug users, based on 2011 National Survey on Drug Use and Health, in Oakland, CA



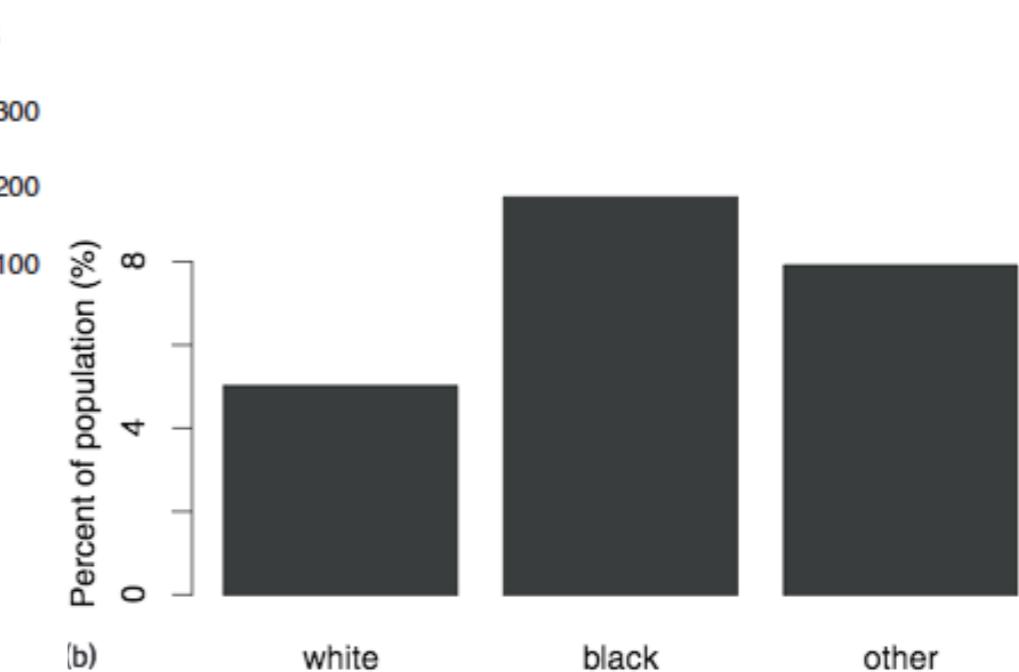
Estimated drug use by race

# Is my data biased? (histograms + geo)

[Lum, Isaac; *Significance*, 2016]

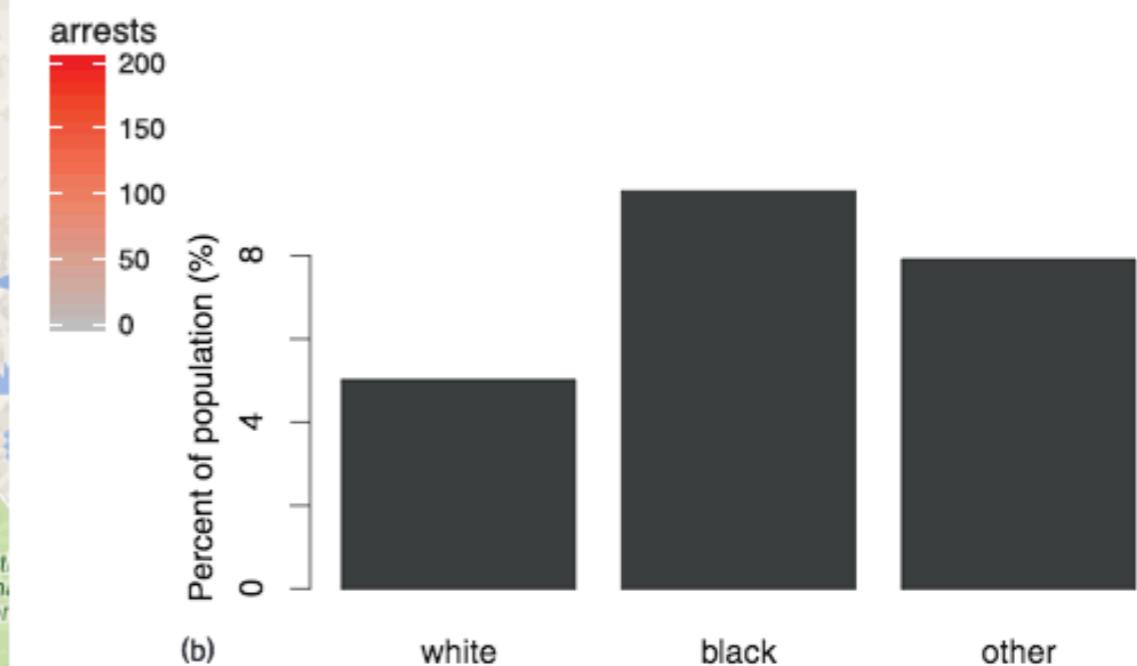
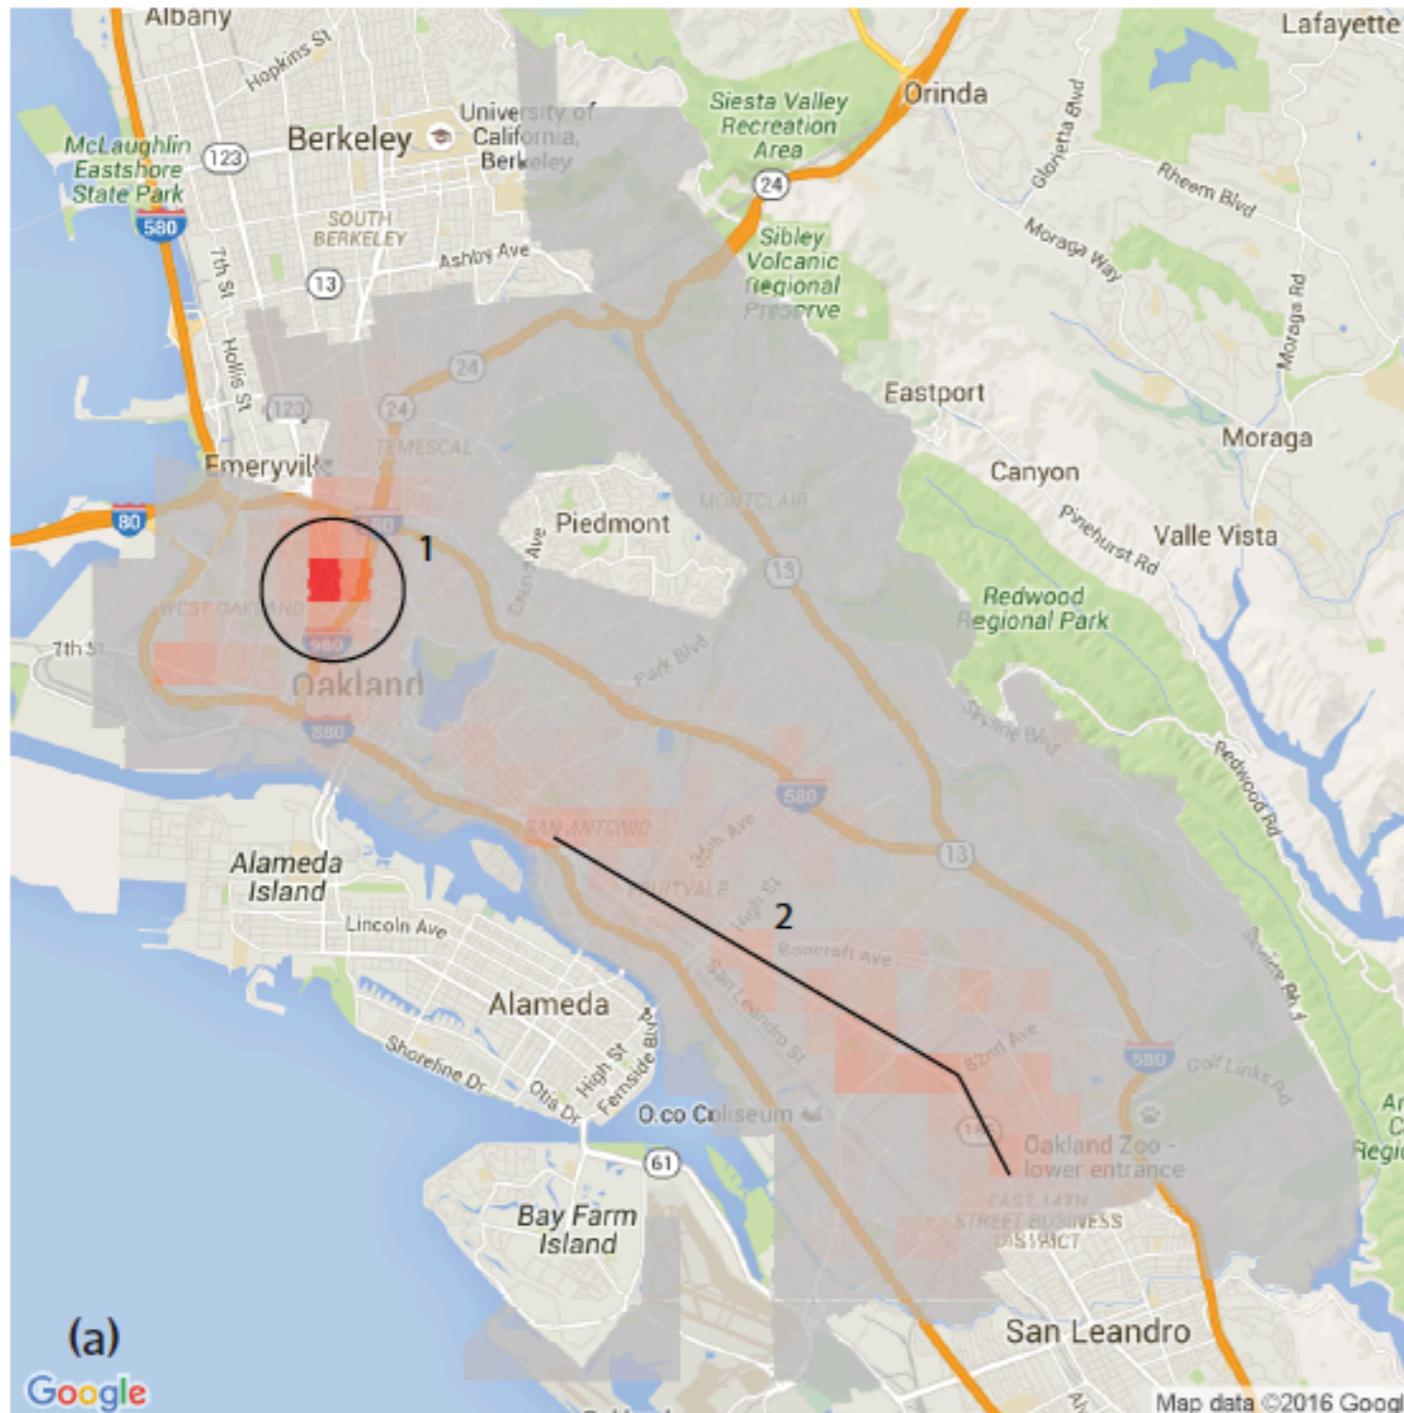


Number of days with targeted policing for drug crimes in areas flagged by PredPol analysis of Oakland, CA, police data for 2011



# Is my data biased? (histograms + geo)

[Lum. Isaac: *Sianificance*. 2016]

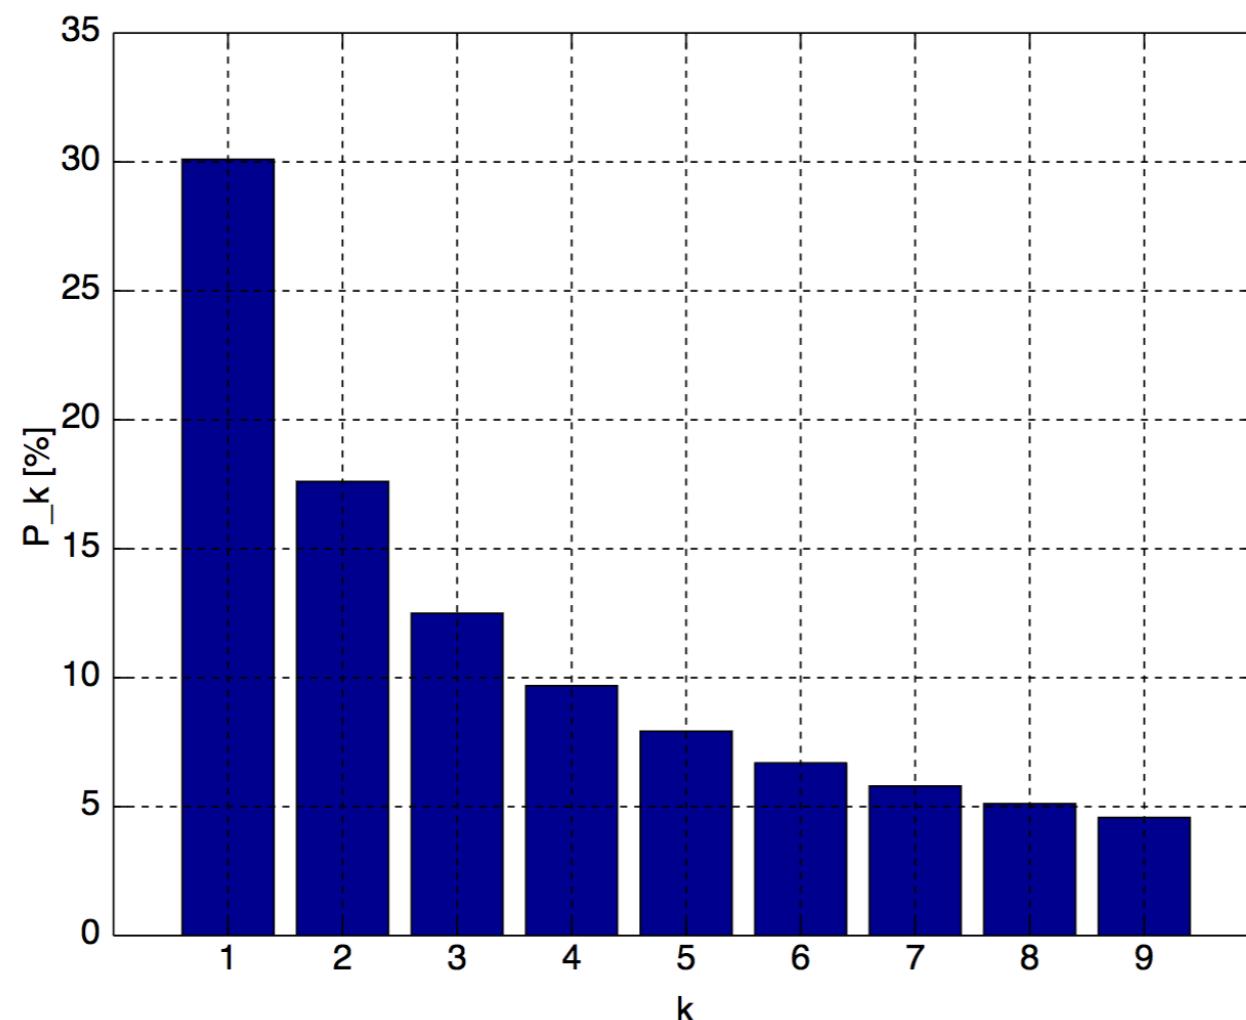


Targeted policing for drug crimes by race

# Benford Law (first digit law)

[Benford: “The law of anomalous numbers” *Proc. Am. Philos. Soc.*, 1938]

The distribution of the first digit **d** of a number, in many naturally occurring domains, approximately follows



$$P(d) = \log_{10} \left( 1 + \frac{1}{d} \right)$$

1 is the most frequent leading digit, followed by 2, etc.

[https://en.wikipedia.org/wiki/Benford%27s\\_law](https://en.wikipedia.org/wiki/Benford%27s_law)

# Benford Law (first digit law)

[Benford: “The law of anomalous numbers” *Proc. Am. Philos. Soc.*, 1938]

The distribution of the first digit **d** of a number, in many naturally occurring domains, approximately follows

$$P(d) = \log_{10} \left( 1 + \frac{1}{d} \right)$$

Holds if  $\log(x)$  is uniformly distributed. **Most accurate** when values are distributed across multiple orders of magnitude, especially **if the process generating the numbers is described by a power law** (common in nature)



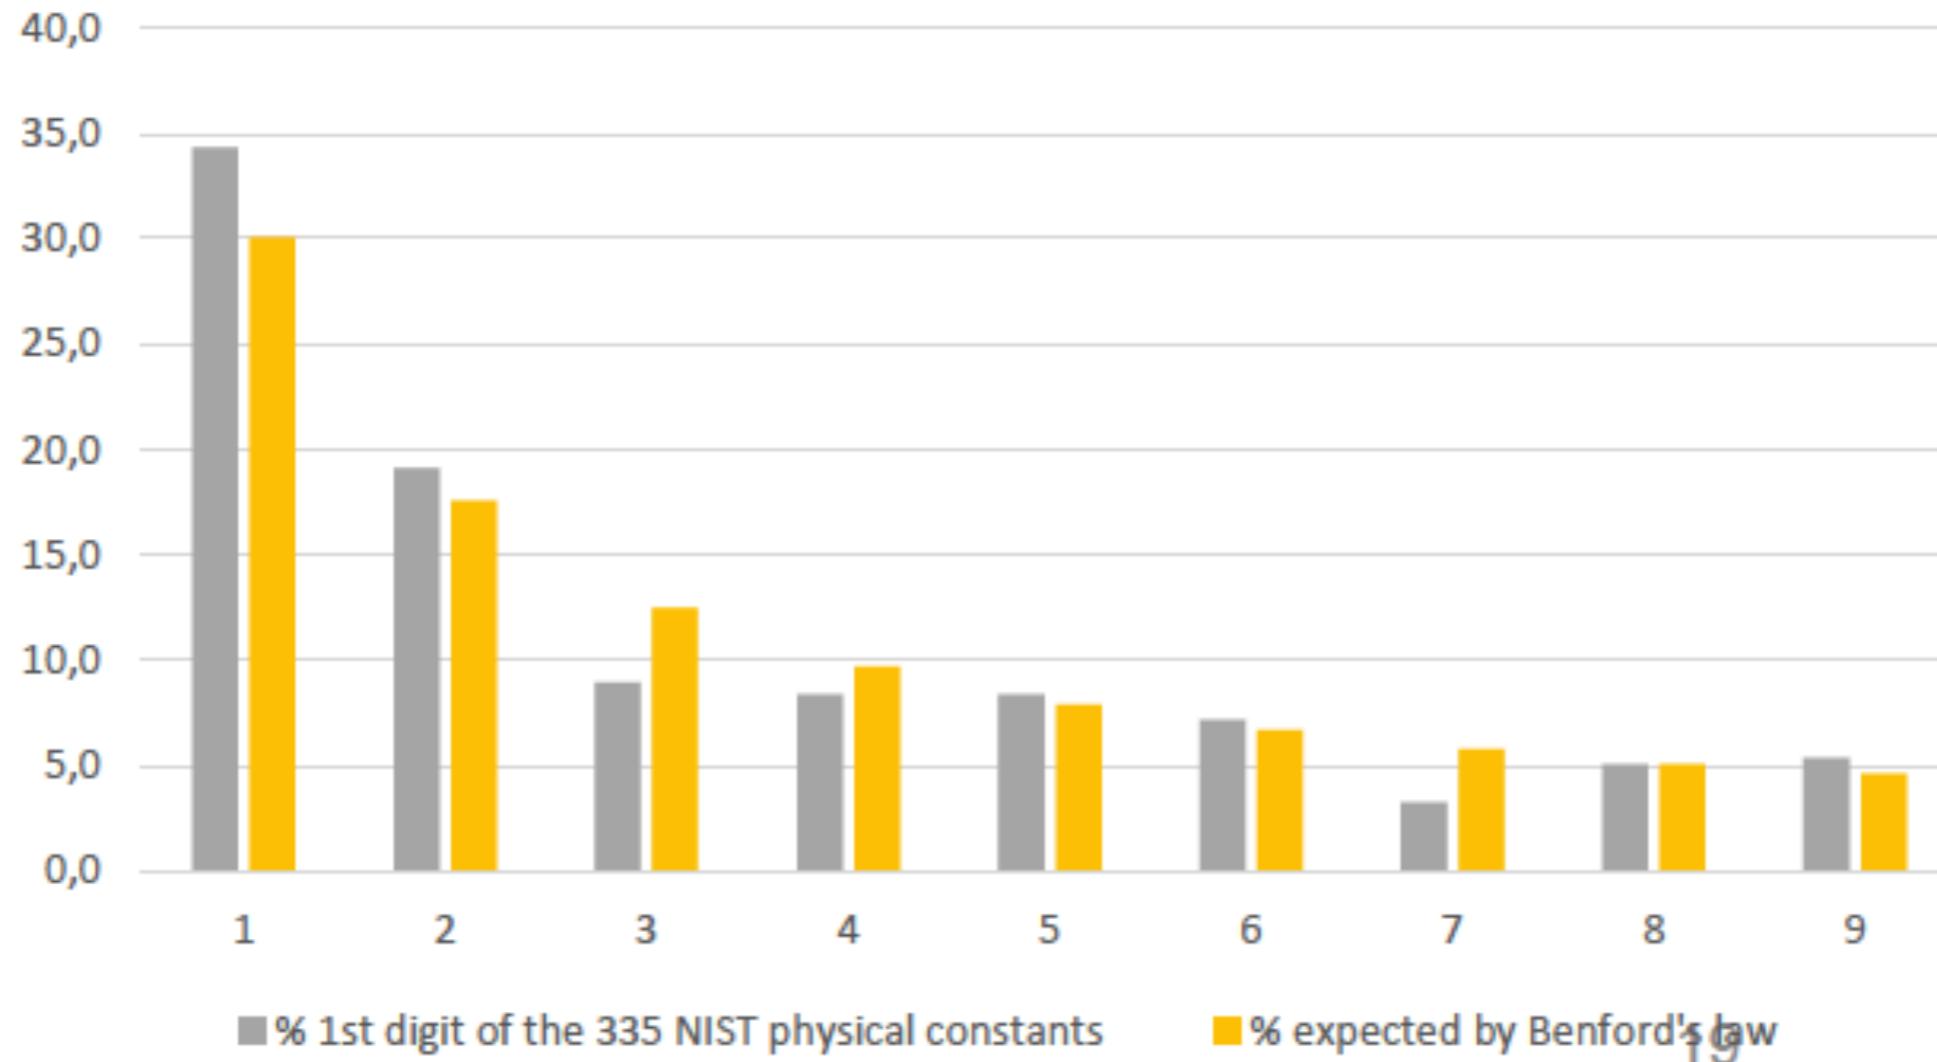
A logarithmic scale bar. Picking a random x position uniformly on this number line, roughly 30% of the time the first digit of the number will be 1.

[https://en.wikipedia.org/wiki/Benford%27s\\_law](https://en.wikipedia.org/wiki/Benford%27s_law)

# Benford Law: an example

[Abedjan, Golab, Naumann; *SIGMOD 2017*]

height of tallest structures



# Benford Law: other examples

[Abedjan, Golab, Naumann; *SIGMOD 2017*]

- surface area of 355 rivers
- sizes of 3,259 US populations
- 104 physical constants
- 1,800 molecular weights
- 308 numbers contained in an issue of Reader's Digest
- Street addresses of the first 342 persons listed in American Men of Science
- ....

**used in fraud detection!**

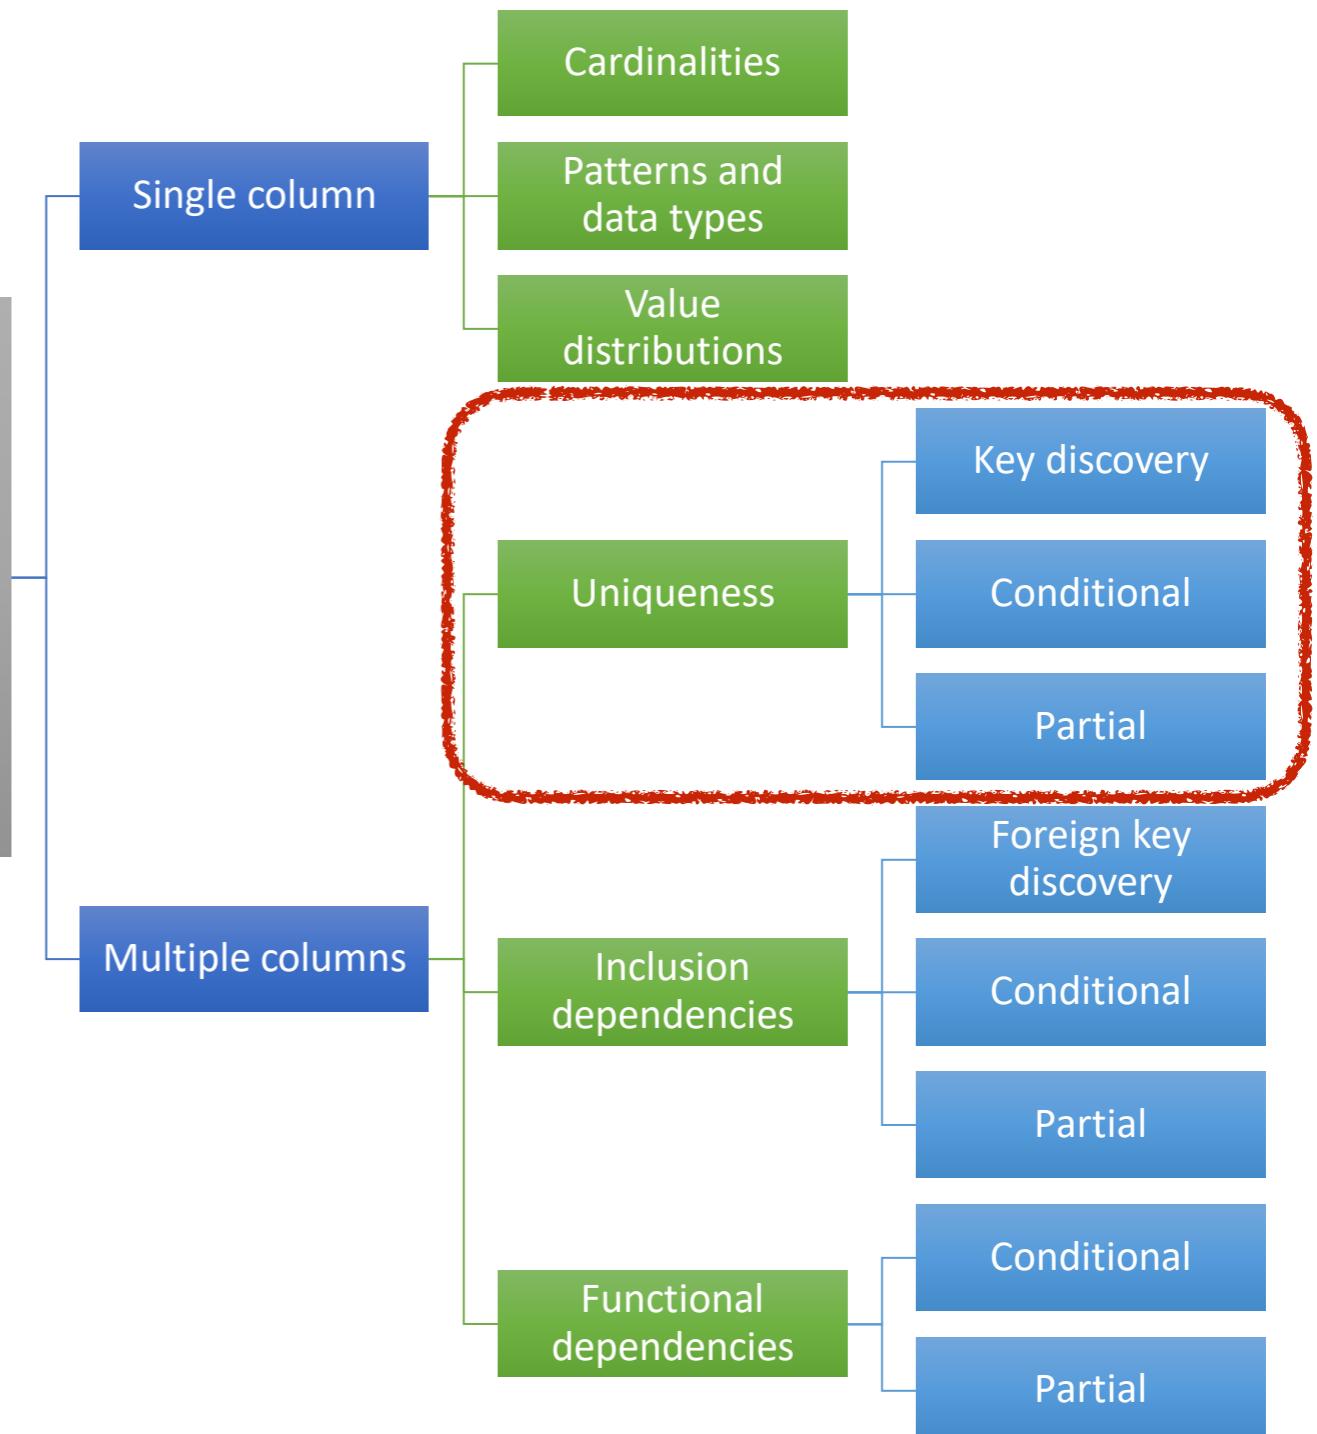
[Benford: “The law of anomalous numbers” *Proc. Am. Philos. Soc.*, 1938]

# Classification of data profiling tasks

[Abedjan, Golab, Naumann; *SIGMOD 2017*]

A	B	C	D	E	F	G	H	
1	UID	sex	race	MarriageSta	DateOfBirth	age	juv_fel_cour	decile_score
2	1	0	1	1	4/18/47	69	0	1
3	2	0	2	1	1/22/82	34	0	3
4	3	0	2	1	5/14/91	24	0	4
5	4	0	2	1	1/21/93	23	0	8
6	5	0	1	2	1/22/73	43	0	1
7	6	0	1	3	8/22/71	44	0	1
8	7	0	3	1	7/23/74	41	0	6
9	8	0	1	2	2/25/73	43	0	4
10	9	0	3	1	6/10/94	21	0	3
11	10	0	3	1	6/1/88	27	0	4
12	11	1	3	2	8/22/78	37	0	1
13	12	0	2	1	12/2/74	41	0	4
14	13	1	3	1	6/14/68	47	0	1
15	14	0	2	1	3/25/85	31	0	3
16	15	0	4	4	1/25/79	37	0	1
17	16	0	2	1	6/22/90	25	0	10
18	17	0	3	1	12/24/84	31	0	5
19	18	0	3	1	1/8/85	31	0	3
20	19	0	2	3	6/28/51	64	0	6
21	20	0	2	1	11/29/94	21	0	9
22	21	0	3	1	8/6/88	27	0	2
23	22	1	3	1	3/22/95	21	0	4
24	23	0	4	1	1/23/92	24	0	4
25	24	0	3	3	1/10/73	43	0	1
26	25	0	1	1	8/24/83	32	0	3
27	26	0	2	1	2/8/89	27	0	3
28	27	1	3	1	9/3/79	36	0	3
29	28	0	3	1	4/17/00	26	0	7

Data profiling



relational data (here: just one table)

# An alternative classification

- To help understand the **statistics**, we look at value ranges, data types, value distributions per column or across columns, etc
- To help understand the **structure** - the (business) rules that generated the data - we look at unique columns / column combinations, dependencies between columns, etc - **reverse-engineer the relational schema** of the data we have
- We need both statistics and structure, they are mutually-reinforcing, and help us understand the **semantics** of the data - it's meaning

**next up: relational model basics**

# The relational model

- Introduced by Edgar F. Codd in 1970 (Turing award)
- At the heart of relational database management systems (RDBMS)
  - a database consists of a collection of **relations** (tables)
  - **tuples** are stored in table rows
  - **attributes** of tuples are stored in table columns

	A	B	C	D	E	F	G	H	
1	UID	sex	race	MarriageSta	DateOfBirth	age	juv_fel	cour	decile_score
2	1	0	1	1	4/18/47	69	0	1	
3	2	0	2	1	1/22/82	34	0	3	
4	3	0	2	1	5/14/91	24	0	4	
5	4	0	2	1	1/21/93	23	0	8	
6	5	0	1	2	1/22/73	43	0	1	
7	6	0	1	3	8/22/71	44	0	1	
8	7	0	3	1	7/23/74	41	0	6	
9	8	0	1	2	2/25/73	43	0	4	
10	9	0	3	1	6/10/94	21	0	3	
11	10	0	3	1	6/1/88	27	0	4	
12	11	1	3	2	8/22/78	37	0	1	
13	12	0	2	1	12/2/74	41	0	4	
14	13	1	3	1	6/14/68	47	0	1	
15	14	0	2	1	3/25/85	31	0	3	
16	15	0	4	4	1/25/79	37	0	1	
17	16	0	2	1	6/22/90	25	0	10	
18	17	0	3	1	12/24/84	31	0	5	
19	18	0	3	1	1/8/85	31	0	3	
20	19	0	2	3	6/28/51	64	0	6	
21	20	0	2	1	11/29/94	21	0	9	
22	21	0	3	1	8/6/88	27	0	2	
23	22	1	3	1	3/22/95	21	0	4	
24	23	0	4	1	1/23/92	24	0	4	
25	24	0	3	3	1/10/73	43	0	1	
26	25	0	1	1	8/24/83	32	0	3	
27	26	0	2	1	2/8/89	27	0	3	
28	27	1	3	1	9/3/79	36	0	3	
29	28	0	2	1	4/27/80	26	0	7	

# The relational model

- Relations are **unordered collections** of tuples
  - conceptually, a relation is a **set** of tuples
  - however, SQL implements a relation as a **multiset** (bag) of tuples
- Why this model?
  - Simple yet powerful. Great for processing very large data sets in bulk

	A	B	C	D	E	F	G	H	
1	UID	sex	race	MarriageSta	DateOfBirth	age	juv_fel	cour	decile_score
2	1	0	1	1	4/18/47	69	0	1	
3	2	0	2	1	1/22/82	34	0	3	
4	3	0	2	1	5/14/91	24	0	4	
5	4	0	2	1	1/21/93	23	0	8	
6	5	0	1	2	1/22/73	43	0	1	
7	6	0	1	3	8/22/71	44	0	1	
8	7	0	3	1	7/23/74	41	0	6	
9	8	0	1	2	2/25/73	43	0	4	
10	9	0	3	1	6/10/94	21	0	3	
11	10	0	3	1	6/1/88	27	0	4	
12	11	1	3	2	8/22/78	37	0	1	
13	12	0	2	1	12/2/74	41	0	4	
14	13	1	3	1	6/14/68	47	0	1	
15	14	0	2	1	3/25/85	31	0	3	
16	15	0	4	4	1/25/79	37	0	1	
17	16	0	2	1	6/22/90	25	0	10	
18	17	0	3	1	12/24/84	31	0	5	
19	18	0	3	1	1/8/85	31	0	3	
20	19	0	2	3	6/28/51	64	0	6	
21	20	0	2	1	11/29/94	21	0	9	
22	21	0	3	1	8/6/88	27	0	2	
23	22	1	3	1	3/22/95	21	0	4	
24	23	0	4	1	1/23/92	24	0	4	
25	24	0	3	3	1/10/73	43	0	1	
26	25	0	1	1	8/24/83	32	0	3	
27	26	0	2	1	2/8/89	27	0	3	
28	27	1	3	1	9/3/79	36	0	3	
29	28	0	2	1	4/27/80	26	0	7	

# The relational model

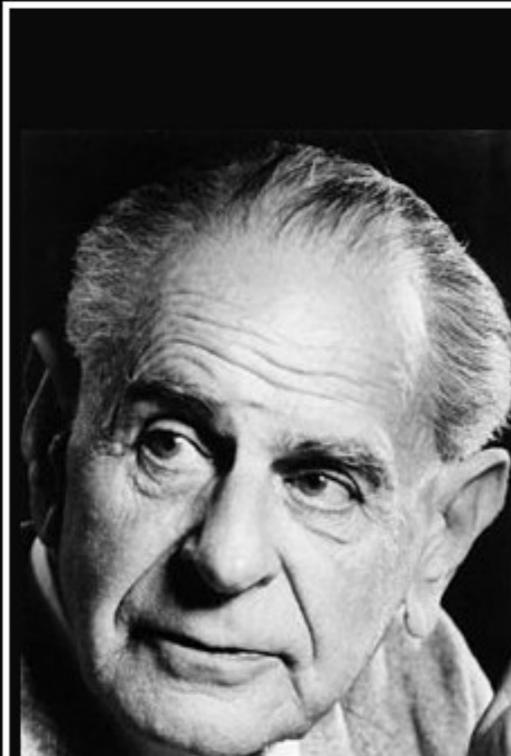
Episodes (season: int, num: int, title: string, viewers: long)

<u>season</u>	<u>num</u>	title	viewers
1	1	Winter is Coming	2.2 M
1	2	The Kingsroad	2.2 M
2	1	The North Remembers	3.9 M
2	2	The Night Lands	3.8 M

- **Relation**: a set or tuples - order doesn't matter, all tuples are distinct
- **Attribute**: a column in a relation (e.g., season)
- **Domain**: data type of an attribute (e.g., season: int)
- **Tuple**: a row in a relation, e.g., (1, 2, The Kingsroad, 2.2 M)

# Schema vs. instances

**Relation schema** is a description of a relation in terms of relation name, attribute names, attribute datatypes, constraints (e.g., keys).  
A schema describes **all valid instances** of a relation.



...no matter how many instances of white swans we may have observed, this does not justify the conclusion that all swans are white.

(Karl Popper)

[izquotes.com](http://izquotes.com)

# Schema vs. instances

**Relation schema** is a description of a relation in terms of relation name, attribute names, attribute datatypes, constraints (e.g., keys). A schema describes **all valid instances** of a relation.

**schema** Episodes (season: integer, num: integer, title: string, viewers: integer)

**instance 1**

<u>season</u>	<u>num</u>	title	viewers
1	1	Winter is Coming	2.2 M
1	2	The Kingsroad	2.2 M
2	1	The North Remembers	3.9 M

**instance 2**

<u>season</u>	<u>num</u>	title	viewers
1	20	Blah, Blah and Blah	0
4	7	Yet Another Title	10 B

**instance 3**

<u>season</u>	<u>num</u>	title	viewers

# Integrity constraints

- Ensure that data adheres to the rules of the application
  - Specified **when schema is defined**
  - Checked and enforced by the DBMS when relations are modified (tuples added / removed / updated)
  - Must **hold on every valid instance** of the database
1. **Domain constraints** - specify valid data types for each attribute, e.g., Students (sid: integer, name: string, gpa: decimal)
  2. **Key constraints** - define a unique identifier for each tuple
  3. **Referential integrity constraints** - specify links between tuples
  4. **Functional dependencies** - show relationships within a table

# Key constraints

A set of attributes is a **candidate key** for a relation if:

- (1) no two distinct tuples can have the same values for all key attributes  
(candidate key **uniquely identifies** a tuple), and
- (2) this is not true for any subset of the key attributes (candidate key **is minimal**)

- If condition (2) is not met, we have a **superkey**
- There may be more than one candidate key for a relation, if so, one is designated as the **primary key**
- All candidate key should be known to property enforce data integrity

**Example:** name possible candidate keys

Students (sid: integer, login: string, name: string, dob: date)

# Key constraints

Example: Students (sid: integer, login: string, name: string, dob: date)

three possible SQL implementations

```
create table Students (
    sid integer      primary key,
    login varchar(128) unique,
    name  varchar(128),
    dob   date,
    gpa   decimal,
    unique (name, dob) );
```

```
create table Students (
    sid integer      unique,
    login varchar(128) primary key,
    name  varchar(128),
    dob   date,
    gpa   decimal,
    unique (name, dob) );
```

```
create table Students (
    sid integer      unique,
    login varchar(128) unique,
    name  varchar(128),
    dob   date,
    gpa   decimal,
    primary key (name, dob) );
```

**NB: every relation must have exactly one primary key**

# DB 101: Where do business rules come from?

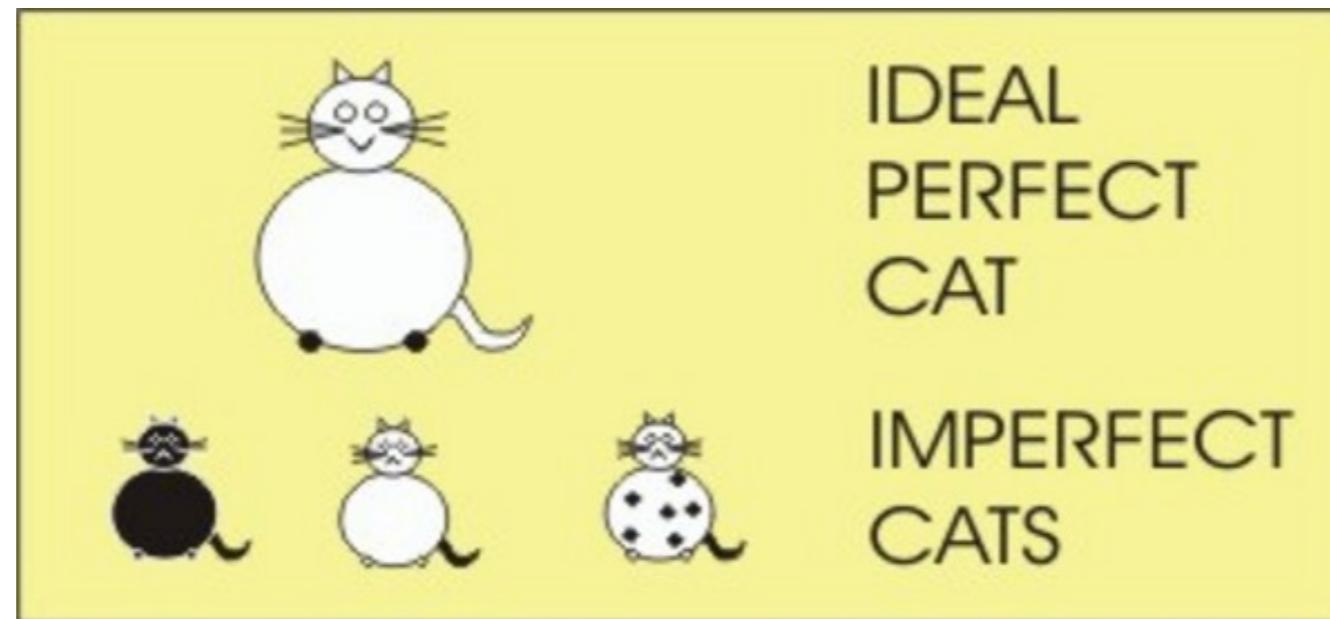
- **Business rules are given**: by the client, by the application designer, by your boss
- Once you know the rules, you create a **relational schema** that encodes these business rules (sometimes starting with the entity-relationship model, sometimes with the relational model directly)
- We can **never-ever-ever deduce business rules by looking at an instance** of a relation!
- We can sometimes know which rules do not hold, but we cannot be sure which rules do hold

Employee

<b>id</b>	<b>login</b>	<b>name</b>
1	jim	Jim Morrison
2	amy	Amy Winehouse
3	amy	Amy Pohler
4	raj	Raj Kapoor

- 1.Which column **is not** a candidate key?
- 2.Which column(s) **may be** a candidate key?
- 3.Give 2 create table statements for which this instance is valid.

# DB (databases) vs. DS (data science)



<https://midnightmediamusings.wordpress.com/2014/07/01/plato-and-the-theory-of-forms/>

- **DB**: start with the schema, admit only data that fits; iterative refinement is possible, and common, but we are still schema-first
- **DS**: start with the data, figure out what schema it fits, or almost fits - reasons of usability, repurposing, low start-up cost
  - the “right” approach is somewhere between these two, data profiling aims to bridge between the two worlds / points of view / methodologies

# Discovering uniques

Given a relation schema  $R$  ( $A, B, C, D$ ) and a relation instance  $r$ , a **unique column combination** (or a “**unique**” for short) is a set of attributes  $X$  whose **projection** contains no duplicates in  $r$

*Episodes(season, num, title, viewers)*

season	num	title	viewers
1	1	Winter is Coming	2.2 M
1	2	The Kingsroad	2.2 M
2	1	The North Remembers	3.9 M

**Projection** is a relational algebra operation that takes as input relation  $R$  and returns a new relation  $R'$  with a subset of the columns of  $R$ .

$\pi_{\text{season}}(\text{Episodes})$

season
1
1
2

$\pi_{\text{season}, \text{num}}(\text{Episodes})$

season	num
1	1
1	2
2	1

$\pi_{\text{title}}(\text{Episodes})$

title
Winter is Coming
The Kingsroad
The North Remembers

# Discovering uniques

Given a relation schema  $R$  ( $A, B, C, D$ ) and a relation instance  $r$ , a **unique column combination** (or a “**unique**” for short) is a set of attributes  $X$  whose **projection** contains no duplicates in  $r$

*Episodes(season, num, title, viewers)*

season	num	title	viewers
1	1	Winter is Coming	2.2 M
1	2	The Kingsroad	2.2 M
2	1	The North Remembers	3.9 M

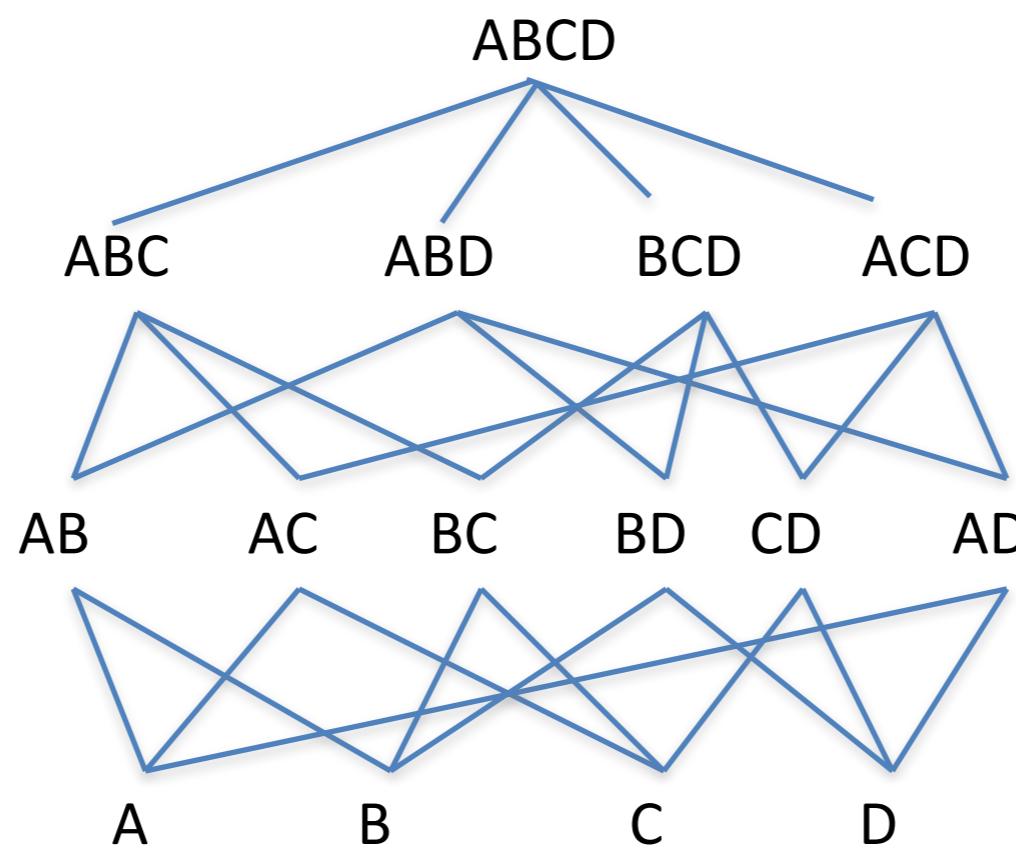
**Projection** is a relational algebra operation that takes as input relation  $R$  and returns a new relation  $R'$  with a subset of the columns of  $R$ .

- Recall that more than one set of attributes  $X$  may be unique
- It may be the case that  $X$  and  $Y$  are both unique, and that they are not disjoint. When is this interesting?

# Discovering uniques

$R (A, B, C, D)$

attribute lattice of  $\mathcal{R}$



$$\begin{pmatrix} 4 \\ 1 \end{pmatrix} = 1$$
$$\begin{pmatrix} 4 \\ 3 \end{pmatrix} = 4$$
$$\begin{pmatrix} 4 \\ 2 \end{pmatrix} = 6$$
$$\begin{pmatrix} 4 \\ 4 \end{pmatrix} = 4$$

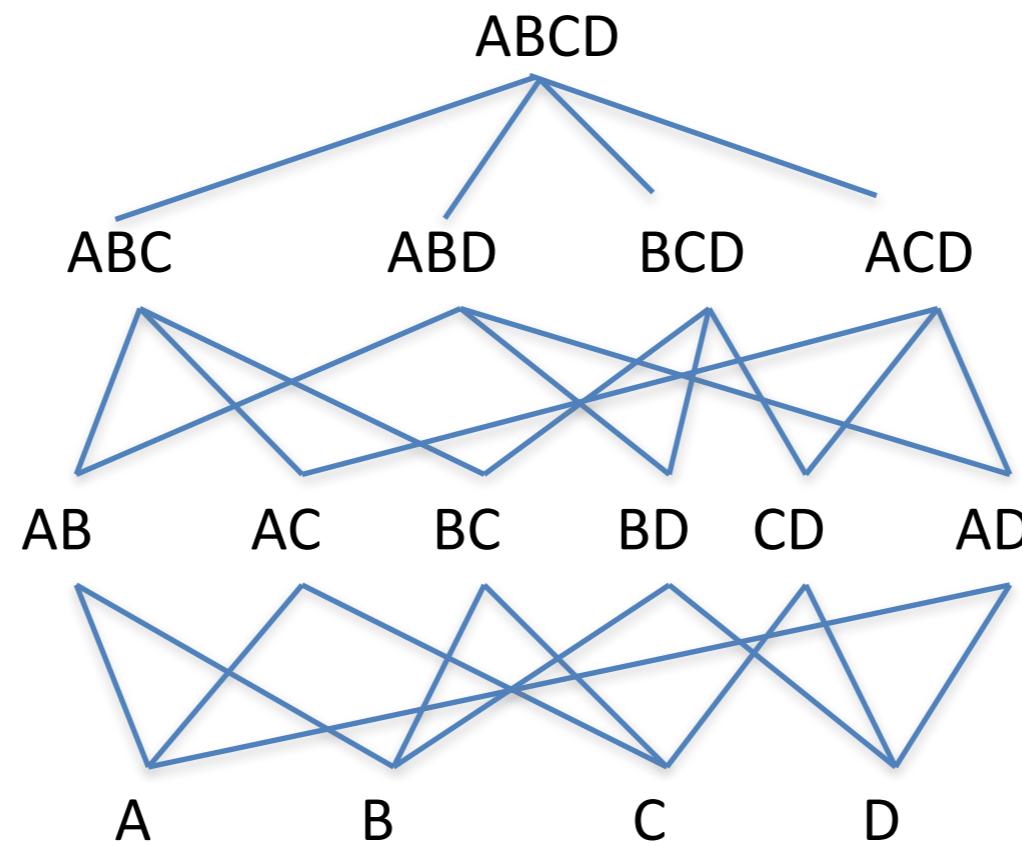
What's the size of the attribute lattice of  $\mathcal{R}$ ?

**Look at all attribute combinations?**

# Discovering uniques

$R (A, B, C, D)$

attribute lattice of  $R$



- If  $X$  is unique, then what can we say about its **superset Y**?
- If  $X$  is non-unique, then what can we say about its **subset Z**?

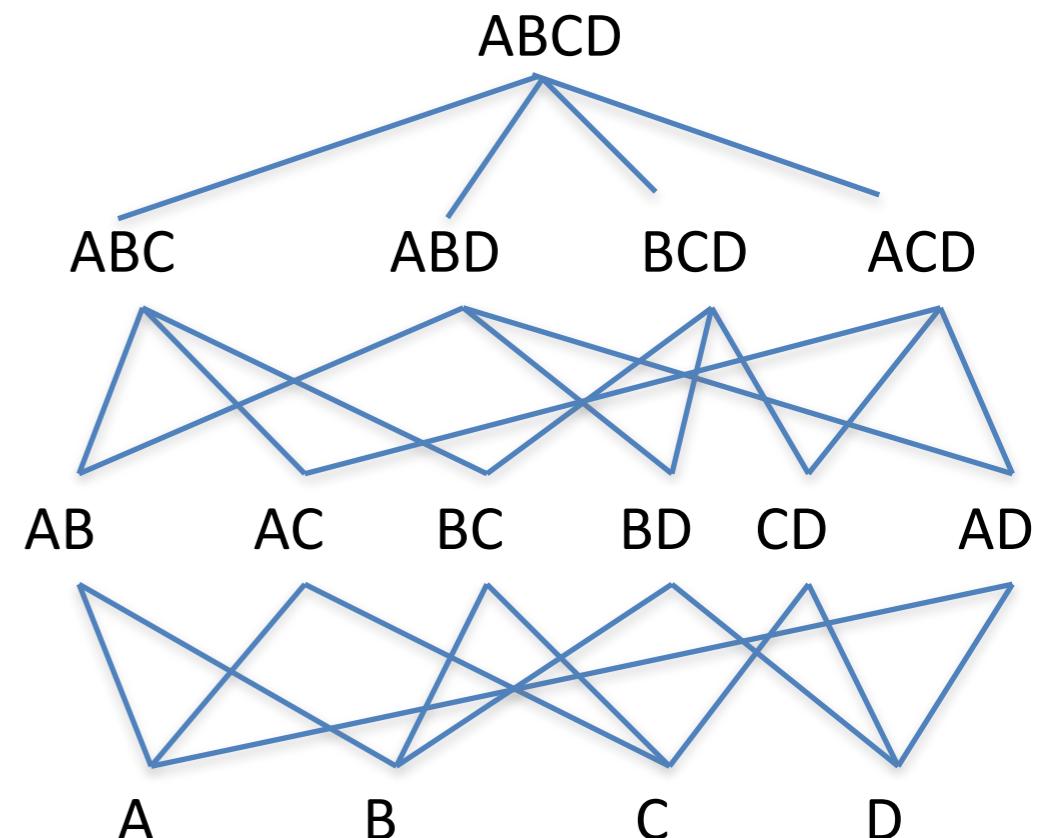
# Discovering uniques

Given a relation schema  $R (A, B, C, D)$  and a relation instance  $r$ , a **unique column combination** (or a “**unique**” for short) is a set of attributes  $X$  whose **projection** contains no duplicates in  $r$

Given a relation schema  $R (A, B, C, D)$  and a relation instance  $r$ , a set of attributes  $Y$  is **non-unique** if its projection contains duplicates in  $r$

$X$  is **minimal unique** if every subset  $Y$  of  $X$  is non-unique

$Y$  is maximal non-unique if every superset  $X$  of  $Y$  is unique



# From uniques to candidate keys

Given a relation schema  $R(A, B, C, D)$  and a relation instance  $r$ , a **unique column combination** is a set of attributes  $X$  whose **projection** contains no duplicates in  $r$

*Episodes(season, num, title, viewers)*

season	num	title	viewers
1	1	Winter is Coming	2.2 M
1	2	The Kingsroad	2.2 M
2	1	The North Remembers	3.9 M

A set of attributes is a **candidate key** for a relation if:

- (1) no two distinct tuples can have the same values for all key attributes (candidate key **uniquely identifies** a tuple), and
- (2) this is not true for any subset of the key attributes (candidate key **is minimal**)

**A minimal unique of a relation instance is a (possible) candidate key of the relation schema.** To find all possible candidate keys, find all minimal uniques in a relation instance.

# Pivot: Frequent itemsets & association rules

- Problem formulation due to Agrawal, Imielinski, Swami, SIGMOD 1993
- Solution: the **Apriori** algorithm by Agrawal & Srikant, VLDB 1994
- Initially for **market-basket data** analysis, has many other applications, we'll see one today
- We wish to answer two related questions:
  - **Frequent itemsets:** Which items are often purchased together, e.g., milk and cookies are often bought together
  - **Association rules:** Which items will likely be purchased, based on other purchased items, e.g., if diapers are bought in a transaction, beer is also likely bought in the same transaction

# Market-basket data

- $I = \{i_1, i_2, \dots, i_m\}$  is the set of available items, e.g., a product catalog of a store
- $X \subseteq I$  is an **itemset**, e.g., {milk, bread, cereal}
- **Transaction**  $t$  is a set of items purchased together,  $t \subseteq I$ , has a transaction id (TID)
  - $t_1$ : {bread, cheese, milk}
  - $t_2$ : {apple, eggs, salt, yogurt}
  - $t_3$ : {biscuit, cheese, eggs, milk}
- Database  $T$  is a set of transactions  $\{t_1, t_2, \dots, t_n\}$
- A transaction  $t$  **supports** an itemset  $X$  if  $X \subseteq t$
- Itemsets supported by at least **minSupp** transactions are called **frequent itemsets**

**minSupp, which can be a number or a percentage, is specified by the user**

# Itemsets

TID	Items
1	A
2	A C
3	A B D
4	A C
5	A B C
6	A B C

***minSupp*** = 2 transactions

How many possible itemsets are there  
(excluding the empty itemset)?

$$2^4 - 1 = 15$$

itemset	support
★ A	6
★ B	3
★ C	4
★ D	1
★ A B	3
★ A C	4
★ A D	1
★ B C	2
★ B D	1
★ C D	0
★ A B C	2
★ A B D	1
★ B C D	0
★ A C D	0
★ A B C D	0

# Association rules

An **association rule** is an implication  $X \rightarrow Y$ , where  $X, Y \subset I$ , and  $X \cap Y = \emptyset$

example:  $\{\text{milk, bread}\} \rightarrow \{\text{cereal}\}$

“A customer who purchased X is also likely to have purchased Y in the same transaction”

we are interested in rules with a **single item** in Y

can we represent  $\{\text{milk, bread}\} \rightarrow \{\text{cereal, cheese}\}$ ?

Rule  $X \rightarrow Y$  holds with **support**  $supp$  in T if  $supp$  of transactions contain  $X \cup Y$

Rule  $X \rightarrow Y$  holds with **confidence**  $conf$  in T if  $conf\%$  of transactions that contain X also contain Y

$$conf \approx \Pr(Y | X)$$

$$conf(X \rightarrow Y) = supp(X \cup Y) / supp(X)$$

# Association rules

**minSupp** = 2 transactions

**minConf** = 0.75

	supp = 3
$A \rightarrow B$	conf = $3 / 6 = 0.5$
$B \rightarrow A$	conf = $3 / 3 = 1.0$

---

	supp = 2
$B \rightarrow C$	conf = $2 / 3 = 0.67$
$C \rightarrow B$	conf = $2 / 4 = 0.5$

---

	supp = 4
$A \rightarrow C$	conf = $4 / 6 = 0.67$
$C \rightarrow A$	conf = $4 / 4 = 1.0$

---

	supp = 2
$AB \rightarrow C$	conf = $2 / 3 = 0.67$
$AC \rightarrow B$	conf = $2 / 4 = 0.5$
$BC \rightarrow A$	conf = $2 / 2 = 1.0$

$$\text{conf}(X \rightarrow Y) = \frac{\text{supp}(X \cup Y)}{\text{supp}(X)}$$

itemset	support
★ A	6
★ B	3
★ C	4
D	1
★ A B	3
★ A C	4
A D	1
★ B C	2
B D	1
CD	0
★ A B C	2
A B D	1
B C D	0
ACD	0
★ A B C D	0

# Association rule mining

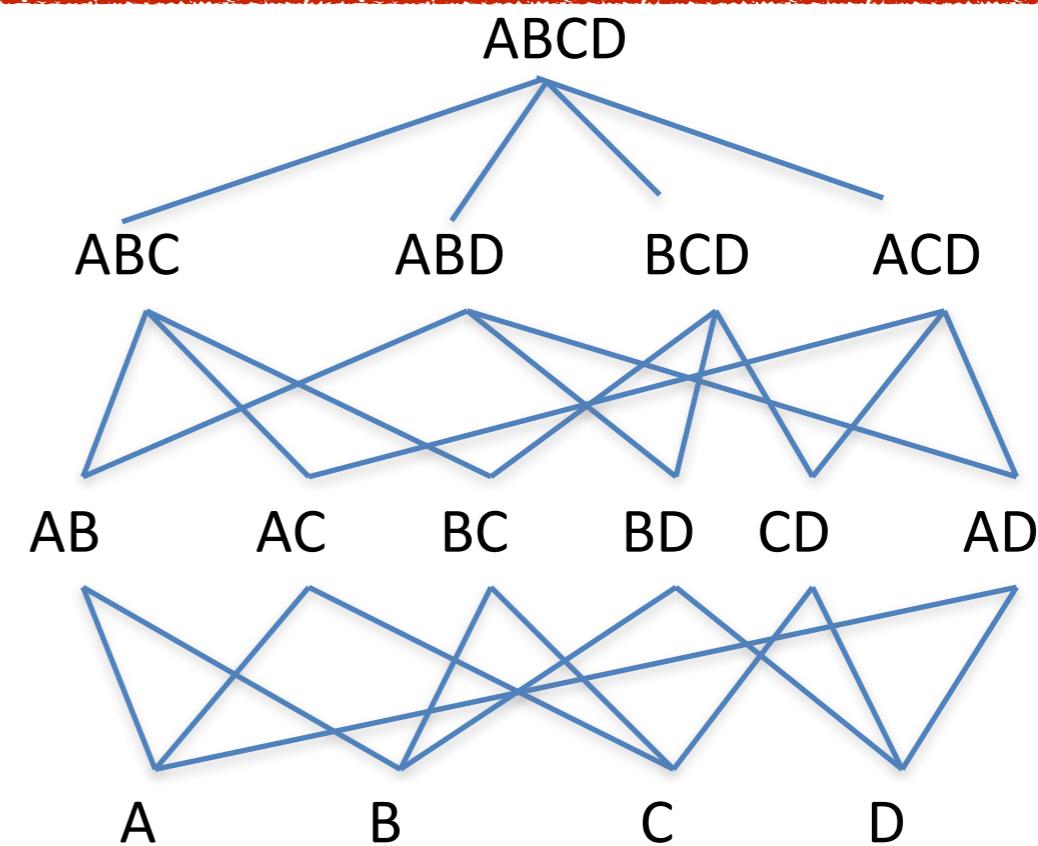
- Goal: find all association rules that satisfy the user-specified minimum support and minimum confidence
- Algorithm outline
  - Step 1: find all frequent itemsets
  - Step 2: find association rules
- Take 1: naïve algorithm for frequent itemset mining
  - Enumerate all subsets of  $I$ , check their support in  $T$
  - **What is the complexity?**

# Key idea: downward closure

itemset	support
A	6
B	3
C	4
D	1
<hr/>	
AB	3
AC	4
AD	1
BC	2
BD	1
CD	0
<hr/>	
ABC	2
ABD	1
BCD	0
ACD	0
<hr/>	
ABCD	0

All subsets of a frequent itemset  $X$  are themselves frequent

So, if some subset of  $X$  is infrequent, then  $X$  cannot be frequent, we know this **apriori**



The converse is not true! If all subsets of  $X$  are frequent,  $X$  is not guaranteed to be frequent

# The Apriori algorithm

**Algorithm Apriori( $T, minSupp$ )**

```
 $F_1 = \{frequent\ 1-itemsets\};$ 
for ( $k = 2; F_{k-1} \neq \emptyset; k++$ ) do
     $C_k \leftarrow \text{candidate-gen}(F_{k-1});$ 
    for each transaction  $t \in T$  do
        for each candidate  $c \in C_k$  do
            if  $c$  is contained in  $t$  then
                 $c.count++;$ 
        end
    end
     $F_k \leftarrow \{c \in C_k \mid c.count \geq minSupp\}$ 
end
return  $F \leftarrow \bigcup_k F_k;$ 
```

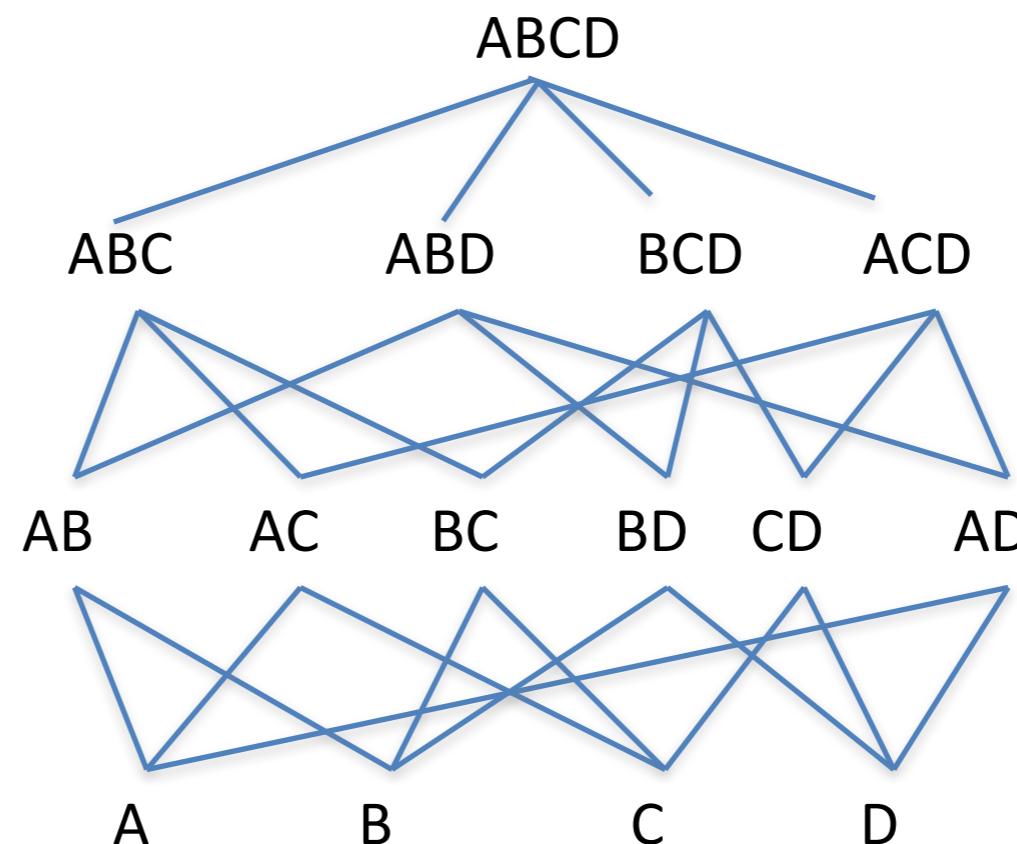
itemset	support
A	6
B	3
C	4
D	1
AB	3
AC	4
AD	1
BC	2
BD	1
CD	0
ABC	2
ABD	1
BCD	0
ACD	0
ABCD	0

# Candidate generation

The **candidate-gen** function takes  $F_{k-1}$  and returns a superset (called the candidates) of the set of all frequent k-itemsets. It has two steps:

**Join:** generate all possible candidate itemsets  $C_k$  of length k

**Prune:** optionally remove those candidates in  $C_k$  that have infrequent subsets



# Candidate generation

Assume a lexicographic ordering of the items

## Join

Insert into  $C_k$   $\langle$

```
select p.item1, p.item2, ..., p.itemk-1, q.itemk-1
from Fk-1 p, Fk-1 q
where p.item1 = q.item1
and p.item2 = q.item2
and ...
and p.itemk-1 < q.itemk-1) why not p.itemk-1 ≠ q.itemk-1?
```

## Prune

```
for each c in Ck do
    for each (k-1) subset s of c do
        if (s not in Fk-1) then
            delete c from Ck
```

# Generating association rules

Rules =  $\emptyset$

**for** each frequent  $k$ -itemset  $X$  **do**

**for** each 1-itemset  $A \subset X$  **do**

        compute conf  $(X / A \rightarrow A) = \text{supp}(X) / \text{sup}(X / A)$

**if** conf  $(X / A \rightarrow A) \geq \text{minConf}$  **then**

$Rules \leftarrow "X / A \rightarrow A"$

**end**

**end**

**end**

**return** Rules

# Performance of *Apriori*

- The possible number of frequent itemsets is exponential,  $O(2^m)$ , where  $m$  is the number of items
- Apriori exploits sparseness and locality of data
  - Still, it may produce a large number of rules: thousands, tens of thousands, ....
  - So, thresholds should be set carefully. **What are some good heuristics?**

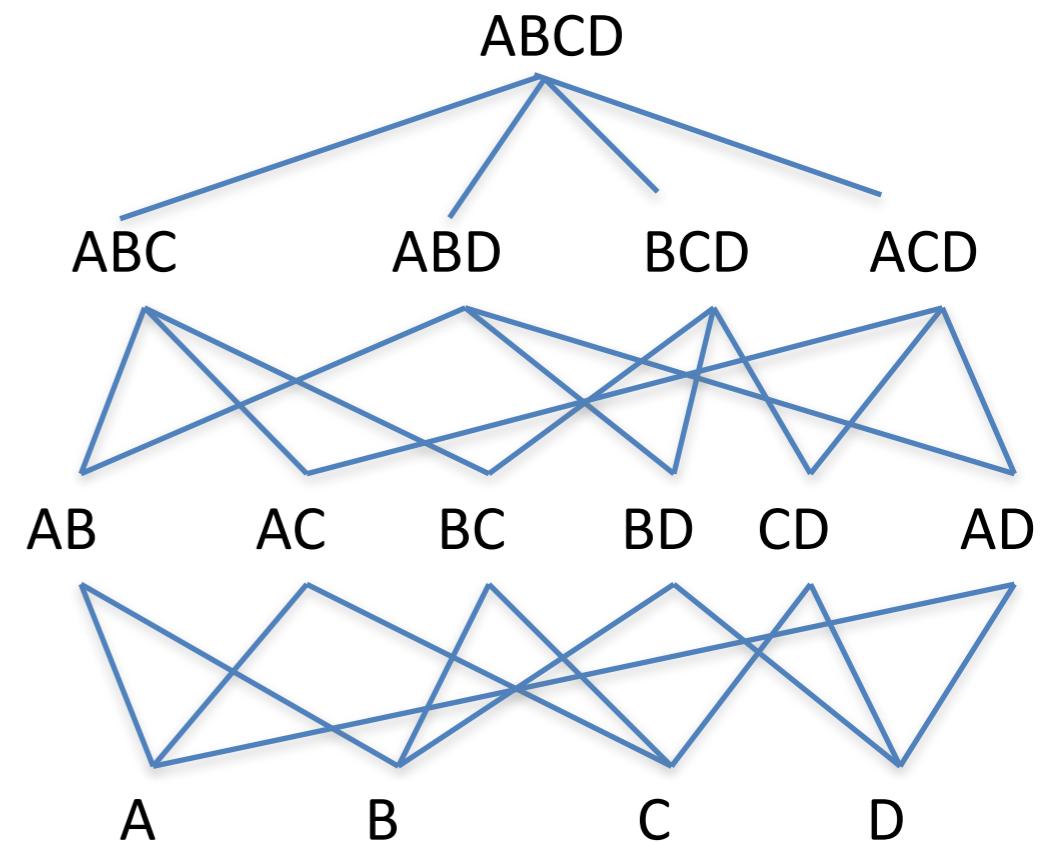
# Back to data profiling: Discovering uniques

Given a relation schema  $R(A, B, C, D)$  and a relation instance  $r$ , a **unique column combination** (or a “**unique**” for short) is a set of attributes  $X$  whose **projection** contains no duplicates in  $r$

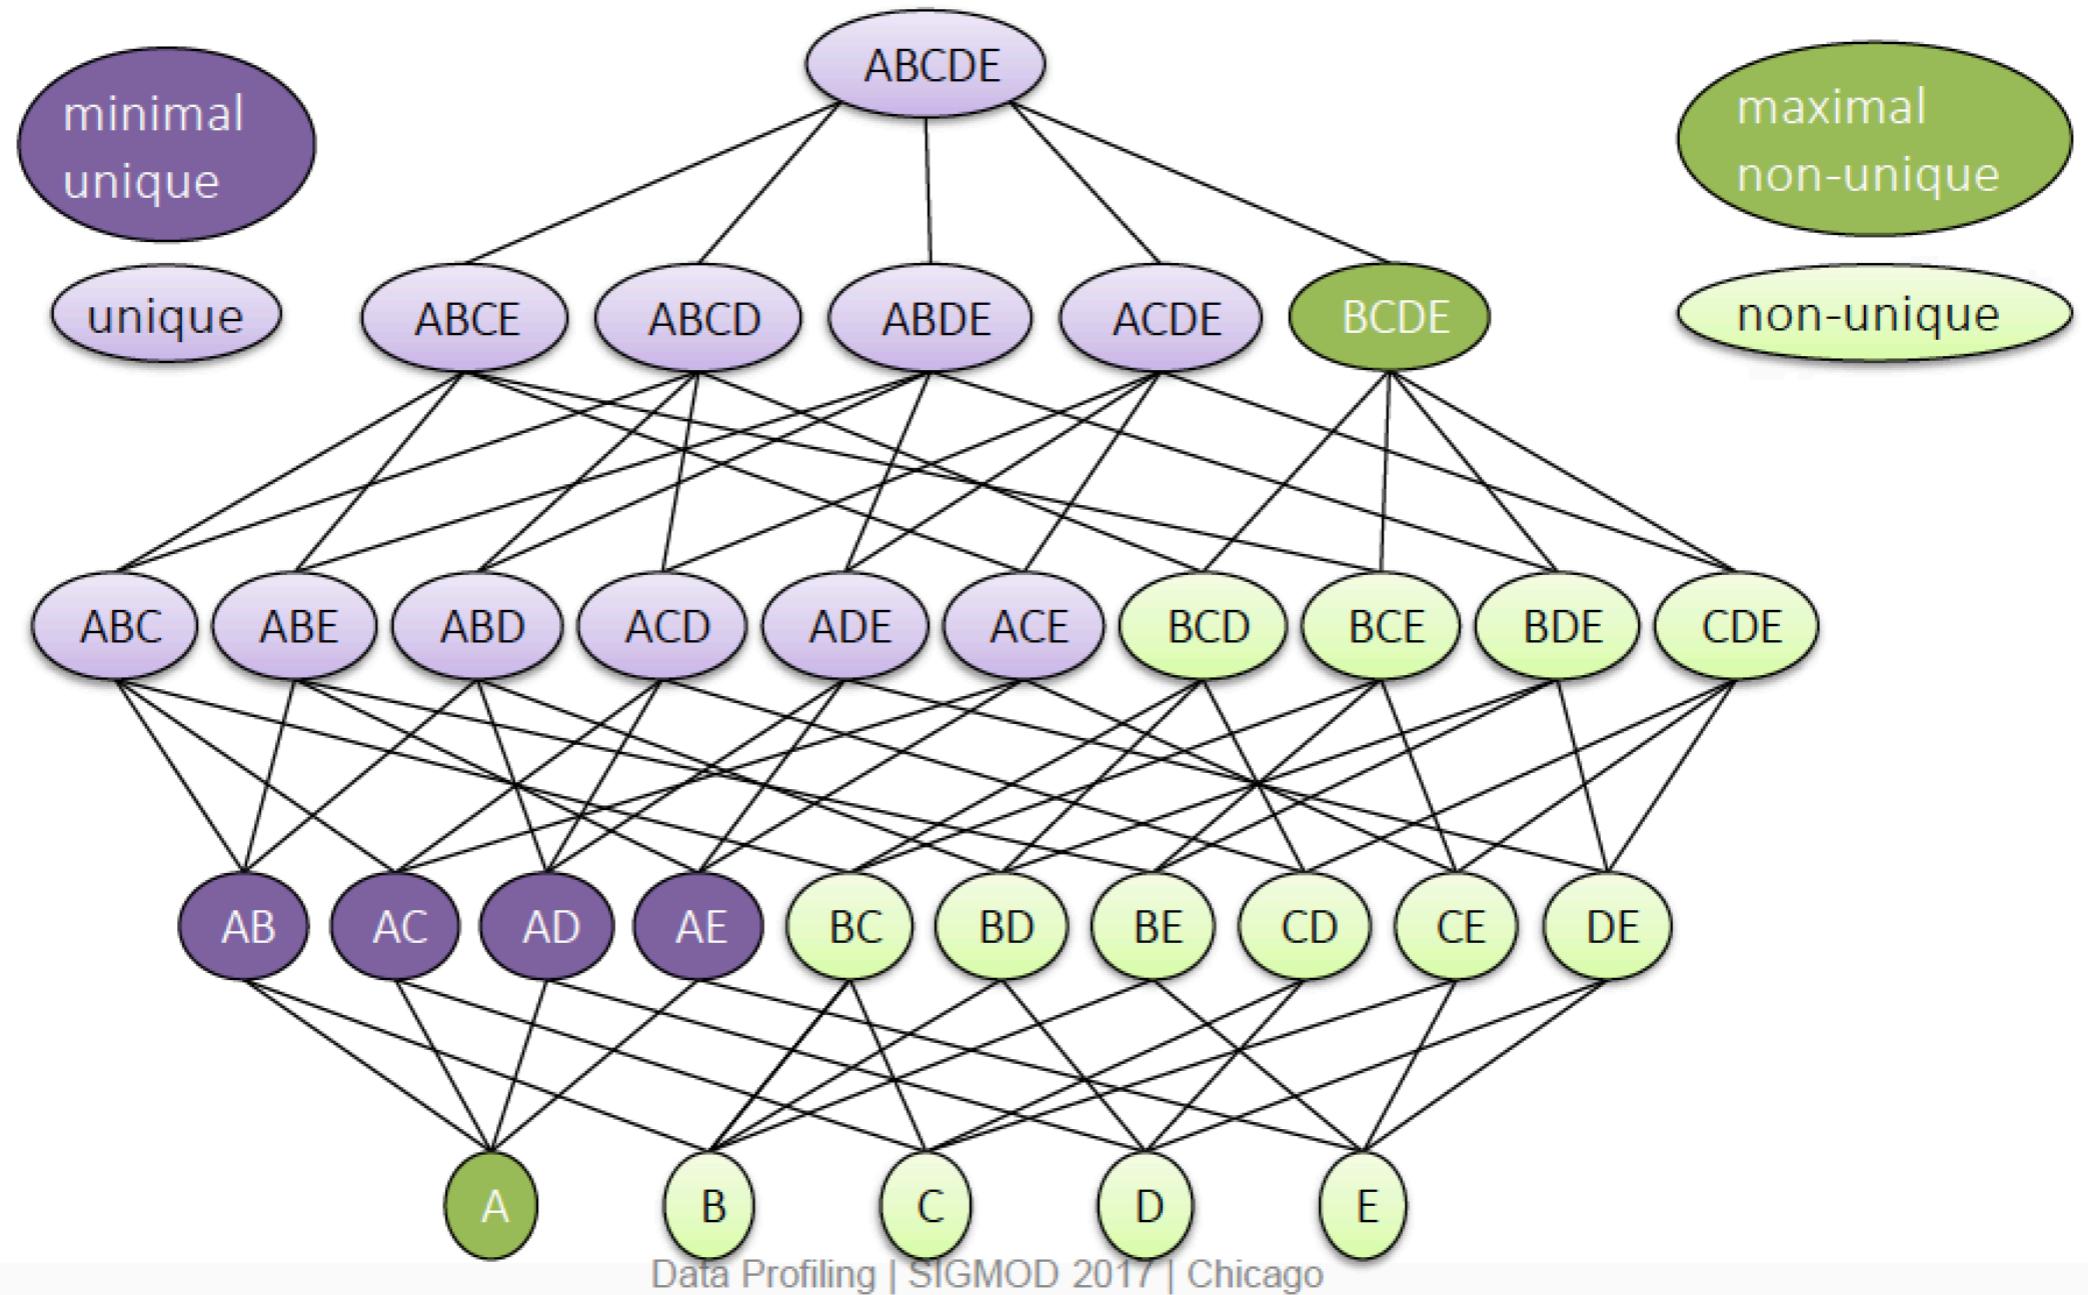
Given a relation schema  $R(A, B, C, D)$  and a relation instance  $r$ , a set of attributes  $Y$  is **non-unique** if its projection contains duplicates in  $r$

$X$  is **minimal unique** if every subset  $Y$  of  $X$  is non-unique

$Y$  is maximal non-unique if every superset  $X$  of  $Y$  is unique



# Output



# From uniques to candidate keys

Given a relation schema  $R (A, B, C, D)$  and a relation instance  $r$ , a **unique column combination** is a set of attributes  $X$  whose **projection** contains no duplicates in  $r$

*Episodes(season, num, title, viewers)*

season	num	title	viewers
1	1	Winter is Coming	2.2 M
1	2	The Kingsroad	2.2 M
2	1	The North Remembers	3.9 M

A set of attributes is a **candidate key** for a relation if:

- (1) no two distinct tuples can have the same values for all key attributes (candidate key **uniquely identifies** a tuple), and
- (2) this is not true for any subset of the key attributes (candidate key **is minimal**)

**A minimal unique of a relation instance is a (possible) candidate key of the relation schema.** To find such possible candidate keys, find all minimal uniques in a given relation instance.

# Apriori-style uniques discovery

[Abedjan, Golab, Naumann; *SIGMOD 2017*]

A **minimal unique** of a relation instance is a **(possible) candidate key** of the relation schema.

**Algorithm Uniques** // sketch, similar to HCA

$U_1 = \{1\text{-uniques}\}$      $N_1 = \{1\text{-non-uniques}\}$

**for** ( $k = 2$ ;  $N_{k-1} \neq \emptyset$ ;  $k++$ ) **do**

$C_k \leftarrow \text{candidate-gen}(N_{k-1})$

$U_k \leftarrow \text{prune-then-check } (C_k)$

// prune candidates with unique sub-sets, and with **value distributions that cannot be unique**

// check each candidate in pruned set for uniqueness

$N_k \leftarrow C_k \setminus U_k$

**end**

breadth-first bottom-up strategy for attribute lattice traversal

return  $U \leftarrow \bigcup_k U_k$ ;